

Envoy Introduction

Suresh Kumar Ponnusamy

April 25, 2018

Outline

- 1 Introduction
- 2 Features
- 3 Components
- 4 Architecture
- 5 Deployment
- 6 Freshdesk + Envoy

What is Envoy I

- L4/L7 proxy and communication bus
- Goal: "The network should be transparent to applications. When network and application problems do occur it should be easy to determine the source of the problem"
- Data plane: Stateless proxy that routes, load balances etc
 - A standalone/stateless proxy with configuration
 - Example: HAProxy, nginx, Envoy, Linkerd etc
- Control plane: Takes the above stateless proxies and turns them into a distributed system.
 - Distributed, dynamic proxy configuration management
 - Example: Istio, Nelson, SmartStack etc

What is Envoy II

- Typical usage: Envoy runs as sidecar process in each machine, all the application communications (ingress/egress) go through this local sidecar Envoy (== "data plane") and Envoy knows how to reach other services. And Envoy itself is configurable dynamically via "control plane"
- Relationship to Istio: Envoy is the "data plane" and other components in Istio fill the "control plane" role, among others.
- More info here

Features I

- Out of process: Can support any language, not a library
- Proxy: L3/L4, L7
- Routing: L7
- Load balancing: with circuit breaking, global rate limiting, request shadowing, outlier detection (== passive health checking) etc
- Edge proxy
 - TLS(+SNI)
 - HTTP 1.1/2
 - L7 routing
- Configuration
 - Static configuration: Can start standalone, pre-configured
 - Dynamic configuration: Almost all aspects of Envoy are configurable on the fly, without having to restart the proxy
- Observability: statsd, distributed tracing support

Features II

- HTTP/2 support
- gRPC support
 - Load balancing
 - HTTP/1.1 bridge
- Filters: L4, L7
 - MongoDB L7 support: sniffing, stats, logging
 - DynamoDB L7 support: sniffing, stats, logging
 - Other custom filters one can write
- Service discovery
 - DNS
 - REST: gRPC based service discovery
- Health checking
- Hot restart

Components/terminology I

- Downstream: The one connecting to Envoy
- Upstream: The one Envoy connects/forwards to
- Listeners: A named network location that downstream clients connect to
 - HAProxy equivalent is "frontend"
 - Dynamically configurable via Listener Discovery Service (LDS)
- Clusters: A group of logically similar upstream hosts that Envoy connects to
 - HAProxy equivalent is "backend"
 - Cluster membership can be statically configured or discovered via Service Discovery
 - DNS
 - Endpoint Discovery Service (EDS), a REST/gRPC based service
 - Active health checking
 - Requests are load balanced among the cluster members based on load balancer policy

- Cluster themselves are discovered via Cluster Discovery Service (CDS)
- Routes: L7 routing
 - HAProxy equivalent is "ACL"
 - Routes can be statically configured or discovered via Route Discovery Service (RDS)

- Envoy is out of process
- Envoy is multi-threaded + evented
 - There are N number of threads (usually equal to number of CPUs, but is configurable).
 - Each thread has its own event loop (using libevent)
 - A connection once accepted stays with the accepting thread, so code doesn't have to be multi-thread aware for most use cases.
- Listeners
 - TCP listeners
 - Listener filters
 - On new connection, filter chain associated with listener is invoked
 - This is where most of the proxy functionality is in: for example, http proxy is a filter chained here
 - Listener Discovery Service (LDS): Service used for adding/removing listeners
- L3/L4 filters

- HTTP connection management filter
 - A filter that implements http protocol: http 1.1, http 2.0, WebSockets
 - access log, requestid generation, distributed tracing, header manipulation, stats etc
 - Route table, configured either statically or dynamically via Route Discovery Service (RDS)
 - HTTP routing
 - Path based (prefix, exact)
 - Header matching
 - WS upgrades
 - Traffic shifting
 - Traffic splitting
 - Retries
 - HTTP filters
- gRPC
 - Bridge filter for languages that has immature gRPC implementation. Client can send requests in http 1.1, Envoy will do the necessary translation.

Design/Architecture III

- gRPC JSON proxy: Client can send requests with JSON payload, Envoy will do the necessary translation and proxy it to gRPC service
- gRPC-Web support: Clients can send a gRPC request to Envoy over HTTP/1.1, Envoy can translate and proxy it to gRPC server
- Envoy uses gRPC in both data plane and control plane
- Cluster manager
 - Cluster Discovery Service (CDS) is used for adding/removing clusters
 - Cluster members discovered via Endpoint Service Discovery (EDS)
- Health check
 - Active and passive
 - Outlier detection (consecutive 5xxs, response time is greater than other machines in the same cluster etc)
- TLS
 - SNI, ALPN, CRL, Session Resumption
- Load balancers
 - Circuit breaking

Design/Architecture IV

- Rate limiting
- Statistics
 - Statsd, DogStatsd, metrics_service (gRPC based one)
- Distributed tracing
 - Request ID generation
 - Integration to external services: LightStep, Zipkin, Jaeger etc
- TCP Proxy
- MongoDB, DynamoDB, Redis filters
- Configuration update
 - Static configuration
 - Hotrestart
 - via DNS (limitations apply)
 - Dynamic configuration
 - For cluster members using EDS
 - For clusters, using CDS
 - For routes, using RDS
 - For listeners, using LDS

Deployment types I

- Sidecar proxy for "service to service"




Figure: Service to Service sidecar

Deployment types II

- Edge proxy




Figure: Edge/Front proxy

- Double edge proxy

Deployment types III




Figure: Double edge proxy

Why Envoy, what value can it add for Freshdesk I

- Better observability within the infra
 - Currently we rely on APM on app side and ELB CloudWatch on edge, everything else in-between is invisible
 - This is going to be a much bigger problem going forward when more microservices are deployed
- Service discovery
 - Currently we rely on OpsWorks to do service discovery
- Clustering
 - Currently we rely on OpsWorks to do the layering, which gets translated into HAProxy "backend" via Chef recipe. Adding/removing machines is clumsy, takes time and error prone.
 - Could potentially help us on minipod effort as well
- Configuration changes
 - Currently we rely on OpsWorks + Stack Settings + Chef recipes to push configuration changes on the fly
 - A system that is designed for dynamic configuration could ease some of these operational challenges

Why Envoy, what value can it add for Freshdesk II

- Customizability
 - We can add custom functionality via filters, both at L7 and L4
 - We can move few things from our Rails app's middleware into Envoy filter, efficient and highly scalable (== apps are not stuck on it, they just do their business logic)
 - Example: when we moved IP blacklisting to HAProxy, we reduced ~10ms on every http request we serve, and improved the overall throughput as well
 - Few things I could think of doing it in Envoy:
 - Ratelimiting
 - WAF features like domain/IP blacklisting etc
 - JWT auth check (already Istio has)?
 - Long shot: MySQL proxying + sharding

Why Envoy, what value can it add for Freshdesk III

- Moving these things out of application, into Envoy, opens up the possibility of choosing appropriate language for each microservice
 - Imagine we want to implement "ticket list" as a microservice in Java
 - If we have most of needed middlewares in Envoy, we can just focus on the business logic (== don't have to worry about how authentication works, how to ratelimit, where/how to find configuration files, how to do sharding etc)
- K8S: Preparing ground for K8S adaption

Possible deployments for Freshdesk I

- As a sidecar in each app machine, in front of nginx/passenger
 - We can configure it to handle all ingress/egress traffic for the application
 - Better observability
 - Other features like circuit breaking, rate limiting etc
 - Distributed tracing?
- As a routing server (replacing HAProxy)
 - Dynamic configuration, service discovery etc
 - We can get out of OpsWorks (w.r.t service discovery and config management)
 - We will still do TLS termination at ELB
- As an edge server
 - Doing TLS termination at our end (for SNI/custom non-wildcard certificates)
 - We will still have NLB/ALB/ELB at the front, just doing TCP termination, firewalling etc.

Q&A