

# Introduction to Linux System Performance Analysis

Suresh Kumar Ponnusamy

September 27, 2017

# Outline

- 1 System
- 2 Application
- 3 Profilers
- 4 Debuggers
- 5 Books

- 1 System
- 2 Application
- 3 Profilers
- 4 Debuggers
- 5 Books

# System performance analysis I

- USE (Utilization Saturation and Errors) Method <sup>1</sup>
- Identify and list the important resources (either physical or software)
  - Example - Hardware: CPU, Memory, Disk, Network etc
  - Example - Software: Locks, processes/threads capacity, file descriptor limit etc
- For each resource
  - Check utilization: How much is it being utilized? Is current utilization "safe"? or Could it lead to problems? It depends on the "resource"
  - Check saturation: Is the resource completely utilized and has extra work queued up? How much is the wait time? Is it within the acceptable range or could it lead to problems?
  - Check errors: Is it malfunctioning? generating some errors?
- A collection of this list + utilities is listed here <sup>2</sup> , <sup>3</sup>, including a picture from there:

# System performance analysis II

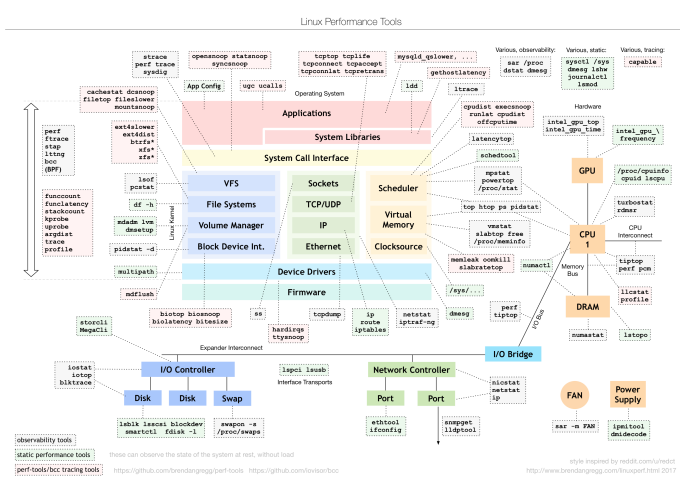


Figure: Linux Performance Tools

<sup>1</sup><http://www.brendangregg.com/usemethod.html>

<sup>2</sup><http://www.brendangregg.com/USEmethod/use-linux.html>

# Resource: CPU I

- uptime/loadavg, top/htop/atop, vmstat, dmstat, "perf sched latency" etc
- Notes
  - loadavg not only shows CPU utilization but also includes other "resources" (i.e., any process that is in UNINTERRUPTIBLE state is also included: primary example is waiting for Disk IO but also includes others like certain locks). For more detailed analysis, you can see here <sup>4</sup>
  - Remember that CPU usage also includes waiting for memory access (Example: trying to access a memory location that is not in cache, it is being brought from RAM, the whole time, CPU will be spinning, waiting for the data)
  - Impact of virtualization: Keep an eye on "steal%"
    - Underlying physical hardware is shared with other workloads from "unknown" individuals/companies
    - Depending on what kind of work it is being done in other instance(s), our performance may be impacted (for example: other instance(s) may be saturating memory bandwidth etc)

# Resource: CPU II

```
# uptime
```

```
uptime
```

```
11:42:54 up 8 days, 1:09, 1 user, load average: 0.24, 0.43, 0.48
```

```
#####
```

```
# Processes that are in running + uninterruptible state
```

```
# ps axl | awk '$10 ~ /[DR]/'
```

```
1    0    79      2  20    0      0    -      R    ?          0:24 [kswapd0]
4    0 24960 24959  20    0 14384 9824 -      D+   pts/22      0:07 dd if=/dev/sda of=/dev/null bs=8M count=10
0    1000 25359 5385  20    0 27636 2800 -      R+   pts/17      0:00 ps axl
```

```
#####
```

```
# vmstat 1
```

```
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
 r  b  swpd   free   buff  cache   si   so    bi   bo    in  cs us sy id wa st
 1  0      0 2002984 748396 8277112    0    0     8   59   33  43  9  4 87  0  0
 0  0      0 1992204 748396 8287912    0    0     0  148 1096 3697  5  1 94  0  0
```

```
#####
```

```
# dstat -c 1
```

```
--total-cpu-usage--
```

```
usr sys idl wai stl
 9  4 87  0  0
 1  1 98  0  0
 3  1 96  0  0
 2  1 98  0  0
```

# Resource: CPU III

```
1 1 98 0 0
```

```
#####
```

```
# top
```

```
top - 11:53:53 up 8 days, 1:20, 1 user, load average: 1.81, 0.96, 0.65
Tasks: 296 total, 1 running, 295 sleeping, 0 stopped, 0 zombie
%Cpu(s): 9.2 us, 3.6 sy, 0.0 ni, 86.9 id, 0.1 wa, 0.0 hi, 0.2 si, 0.0 st
GiB Mem : 15.555 total, 1.843 free, 5.044 used, 8.668 buff/cache
GiB Swap: 0.000 total, 0.000 free, 0.000 used. 8.062 avail Mem
```

PID	USER	PR	NI	VIRT	RES	%CPU	%MEM	TIME+	S	COMMAND
27648	suresh	20	0	1463.9m	514.3m	31.2	3.2	8:56.20	S	chromium

```
#####
```

```
# htop
```

1	[	3.4%	5	[
2	[	2.7%	6	[
3	[	3.3%	7	[
4	[	6.0%	8	[
Mem	[	7.75G/15.6G	Tasks: 188; 2 running	
Swp	[	OK/OK	Load average: 1.42 0.96 0.67	
			Uptime: 8 days, 01:21:02	

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
27648	suresh	20	0	1463M	514M	410M	S	29.4	3.2	9:07.72	/usr/lib/chromium/chromium --type=renderer --fie

```
#####
```

```
# sudo perf sched record sleep 10
```



# Resource: CPU IV

```
# sudo perf sched latency
```

Task	Runtime ms	Switches	Average delay ms	Maximum delay ms	Maximum delay at
kworker/4:1H:155	0.050 ms	1	avg: 0.051 ms	max: 0.051 ms	max at: 262585.22959
khugepaged:67	2.600 ms	1	avg: 0.042 ms	max: 0.042 ms	max at: 262585.44290
ksoftirqd/0:7	0.116 ms	9	avg: 0.042 ms	max: 0.312 ms	max at: 262587.32127
tr: (20)	14.232 ms	20	avg: 0.031 ms	max: 0.387 ms	max at: 262580.03030
kworker/u16:0:31591	1.332 ms	22	avg: 0.028 ms	max: 0.055 ms	max at: 262580.10994
sed: (20)	55.704 ms	22	avg: 0.025 ms	max: 0.211 ms	max at: 262584.03385
perf:18602	1.494 ms	1	avg: 0.025 ms	max: 0.025 ms	max at: 262587.63840
wpa_supplicant:673	0.036 ms	1	avg: 0.024 ms	max: 0.024 ms	max at: 262583.05633
WorkerPool/1344:15512	0.085 ms	1	avg: 0.024 ms	max: 0.024 ms	max at: 262580.73229
WorkerPool/875:28605	2.886 ms	1	avg: 0.023 ms	max: 0.023 ms	max at: 262583.96119
WorkerPool/1553:6841	0.266 ms	1	avg: 0.021 ms	max: 0.021 ms	max at: 262580.45587
WorkerPool/1310:536	0.135 ms	1	avg: 0.021 ms	max: 0.021 ms	max at: 262583.93921
WorkerPool/696:14642	0.169 ms	2	avg: 0.021 ms	max: 0.021 ms	max at: 262579.92471
rcu_preempt:8	12.843 ms	483	avg: 0.021 ms	max: 0.054 ms	max at: 262581.42652
WorkerPool/992:9409	0.163 ms	2	avg: 0.021 ms	max: 0.023 ms	max at: 262580.72839
mozStorage #3:1398	21.238 ms	2	avg: 0.020 ms	max: 0.021 ms	max at: 262581.18988
tmux: (10)	12.865 ms	10	avg: 0.019 ms	max: 0.128 ms	max at: 262578.03307
Chrome_DBThread:10992	0.446 ms	6	avg: 0.019 ms	max: 0.033 ms	max at: 262581.21913

<sup>4</sup><http://brendangregg.com/blog/2017-08-08/linux-load-averages.html>



# Resource: Memory I

- top/htop/atop, smem, vmstat, dmstat etc
- Notes on certain terminology
  - Virtual memory (VSS): Address space used, not an indicative of physical memory usage
  - Resident memory (RSS): Actual physical memory used, including memory shared with other processes. Two kinds of sharing could happen:
    - By forking a process (so both parent and child share same memory), generally with CoW (Copy-on-Write) semantics.
    - By memory mapping same file
  - Unique Set Size (USS): Actual private physical memory used i.e., not including the memory shared with other processes
  - Proportional Set Size (PSS): Private physical memory used + proportion of shared memory with other processes
  - Memory overcommit <sup>5</sup> : Assignment of more memory than physical memory available, assuming not everyone will need all this memory at the same time. For example, Redis needs it:

# Resource: Memory II

- Redis bgsave needs it, since bgsave forks a new process out of existing redis process (which is likely using lots of memory), the new forked process also will "appear" to use more memory but in reality it won't need more memory, it just does bgsave and exits.
- So if we disable overcommit, then Redis fork will fail etc.

Process A has 50 KiB of unshared memory

Process B has 300 KiB of unshared memory

Both process A and process B have 100 KiB of the same shared memory region

RSS of process A = 50KiB + 100KiB = 150 KiB

USS of process A = 50 KiB

PSS of process A = 50 KiB + (100 KiB / 2) = 100 KiB

RSS of process B = 300KiB + 100KiB = 400 KiB

USS of process B = 300 KiB

PSS of process B = 300 KiB + (100 KiB / 2) = 350 KiB

Some example command invocations:

# Resource: Memory III

```
# Total memory
```

```
# free -m
```

	total	used	free	shared	buffers	cached
Mem:	7482	5479	2002	0	154	1573
-/+ buffers/cache:		3751	3731			
Swap:	0	0	0			

```
#####
```

```
# dstat -m 1
```

```
-----memory-usage-----
used free buff cach
8041M 1729M 734M 7682M
8022M 1748M 734M 7662M
8021M 1749M 734M 7662M
8022M 1748M 734M 7661M
8021M 1749M 734M 7661M
```

```
#####
```

```
# vmstat 1
```

```
procs -----memory----- --swap-- ----io---- -system-- -----cpu-----
r  b  swpd  free  buff  cache  si  so  bi  bo  in  cs  us  sy  id  wa  st
1  0      0 1834808 751556 8327776  0  0   8  57  36  53  9  4  87  0  0
0  0      0 1835048 751556 8327772  0  0   0  0  685 1479 1  1  98  0  0
0  0      0 1835184 751556 8327516  0  0   0  0  756 1770 1  1  98  0  0
0  0      0 1835616 751556 8327512  0  0   0  0  669 1710 1  1  98  0  0
```

```
#####
```

# Resource: Memory IV

```
# top
top - 11:53:53 up 8 days, 1:20, 1 user, load average: 1.81, 0.96, 0.65
Tasks: 296 total, 1 running, 295 sleeping, 0 stopped, 0 zombie
%Cpu(s): 9.2 us, 3.6 sy, 0.0 ni, 86.9 id, 0.1 wa, 0.0 hi, 0.2 si, 0.0 st
GiB Mem : 15.555 total, 1.843 free, 5.044 used, 8.668 buff/cache
GiB Swap: 0.000 total, 0.000 free, 0.000 used. 8.062 avail Mem
```

PID	USER	PR	NI	VIRT	RES	%CPU	%MEM	TIME+	S	COMMAND
27648	suresh	20	0	1463.9m	514.3m	31.2	3.2	8:56.20	S	chromium

#####

```
# htop --sort-key=RES
```

1	[	0.0%	5	[
2	[	0.0%	6	[
3	[	0.0%	7	[
4	[	0.0%	8	[

Mem[|||||]7.94G/15.6G Tasks: 190; 1 running  
Swp[ OK/OK] Load average: 0.48 0.53 0.35  
Uptime: 8 days, 02:00:49

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
1223	suresh	20	0	3667M	1334M	267M	S	0.0	8.4	36:59.01	/usr/lib/firefox/firefox https://docs.google.com
11054	suresh	20	0	1156M	517M	420M	S	0.0	3.2	1:37.12	/usr/lib/chromium/chromium --type=gpu-process --
16413	suresh	20	0	1052M	470M	27264	S	0.0	3.0	7:54.46	./src/emacs/src/emacs
12621	suresh	20	0	1439M	459M	273M	S	0.0	2.9	0:49.31	/usr/lib/chromium/chromium --type=renderer --fie

#####

# Resource: Memory V

```
# smem -k -t -w
```

Area	Used	Cache	Noncache
firmware/hardware	0	0	0
kernel image	0	0	0
kernel dynamic memory	1.9G	1.8G	88.2M
userspace memory	3.7G	36.4M	3.7G
free memory	1.7G	1.7G	0
-----			
	7.3G	3.5G	3.8G

```
# User-wise memory usage
```

```
smem -t -k -u
```

User	Count	Swap	USS	PSS	RSS
rpc	1	0	404.0K	440.0K	2.0M
dbus	1	0	552.0K	586.0K	2.0M
nagios	1	0	692.0K	748.0K	3.1M
rpcuser	1	0	800.0K	862.0K	3.0M
ntp	1	0	764.0K	922.0K	3.9M
smmsp	1	0	1.4M	1.6M	3.4M
nobody	1	0	1.1M	1.9M	7.7M
suresh	2	0	936.0K	2.4M	8.4M
aws	4	0	239.6M	275.3M	394.1M
root	32	0	279.9M	286.1M	353.0M
deploy	19	0	2.9G	3.2G	4.6G
-----					
	64	0	3.4G	3.8G	5.4G

# Resource: Memory VI

# Application-wise memory usage

# smem -k -t

PID	User	Command	Swap	USS	PSS	RSS
2971	root	/sbin/minigetty /dev/tty5	0	84.0K	106.0K	1.3M
2966	root	/sbin/minigetty /dev/tty3	0	88.0K	110.0K	1.4M
2973	root	/sbin/minigetty /dev/tty6	0	88.0K	110.0K	1.4M
.....						
21207	root	Passenger core	0	6.9M	7.7M	12.8M
18816	root	python /usr/bin/smem -k -t	0	8.1M	8.5M	10.8M
7154	root	/usr/bin/python /usr/bin/le	0	20.8M	21.2M	24.6M
2991	aws	opsworks-agent: master 2991	0	18.9M	27.8M	56.5M
3013	aws	opsworks-agent: statistics	0	52.1M	61.1M	91.2M
2994	aws	opsworks-agent: keep_alive	0	61.6M	70.6M	100.6M
3017	aws	opsworks-agent: process_com	0	107.0M	115.8M	145.8M
16816	root	/opt/SumoCollector/jre/bin/	0	228.0M	228.1M	230.3M
21253	deploy	Passenger AppPreloader: /da	0	187.8M	246.8M	474.7M
15144	deploy	Passenger RubyApp: /data/he	0	220.3M	278.4M	501.2M
7714	deploy	Passenger RubyApp: /data/he	0	281.2M	322.2M	530.4M
23074	deploy	Passenger RubyApp: /data/he	0	328.9M	364.1M	561.7M
20298	deploy	Passenger RubyApp: /data/he	0	531.9M	562.7M	746.9M
20926	deploy	Passenger RubyApp: /data/he	0	658.8M	690.0M	871.8M
16513	deploy	Passenger RubyApp: /data/he	0	753.2M	783.0M	957.8M
-----						
64	11		0	3.4G	3.8G	5.4G

# Memory usage by mapping

# smem -k -t -m

Map	PIDs	AVGPSS	PSS
/[aio]	12	0	0
/data/helpkit/shared/bundler_gems/ruby/2	7	0	0

# Resource: Memory VII

```
/opt/SumoCollector/19.182-44/lib/aether-      1      0      0
/opt/SumoCollector/19.182-44/lib/aether-      1      0      0
/opt/SumoCollector/19.182-44/lib/akka-ac      1      0      0
.....
/usr/sbin/nginx                             13    58.0K    755.0K
/usr/lib64/perl5/CORE/libperl.so             13    63.0K    823.0K
/usr/lib64/libssl.so.1.0.1k                  32    27.0K    895.0K
/bin/bash                                    2    448.0K    896.0K
/usr/lib64/libnss3.so                        15    80.0K     1.2M
/usr/lib64/libkrb5.so.3.3                    39    33.0K     1.3M
/usr/lib64/libxml2.so.2.9.1                  25    54.0K     1.3M
/usr/lib64/libpython2.7.so.1.0               2    890.0K     1.7M
/usr/local/lib/ruby/gems/2.2.0/bundler/g      3    680.0K     2.0M
/opt/aws/opsworks/local/bin/ruby             4    556.0K     2.2M
[stack]                                     64    35.0K     2.2M
/usr/local/lib/libruby.so.2.2.0              7    371.0K     2.5M
/lib64/libc-2.17.so                         64    41.0K     2.6M
/lib64/libcrypto.so.1.0.1k                  36   101.0K     3.6M
/opt/SumoCollector/jre/lib/amd64/server/      1     6.8M     6.8M
<anonymous>                                64     4.1M    260.1M
[heap]                                       64   45.7M     2.9G
-----
448                                           2956    66.5M    3.2G
```

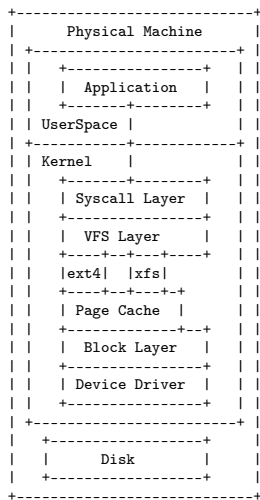
```
# See here for possible values: https://www.kernel.org/doc/Documentation/vm/overcommit-accounting
# sysctl vm.overcommit_memory
vm.overcommit_memory = 0
```

```
# Out of memory errors
dmesg -T | grep OOM
```

<sup>5</sup><https://www.etalabs.net/overcommit.html>

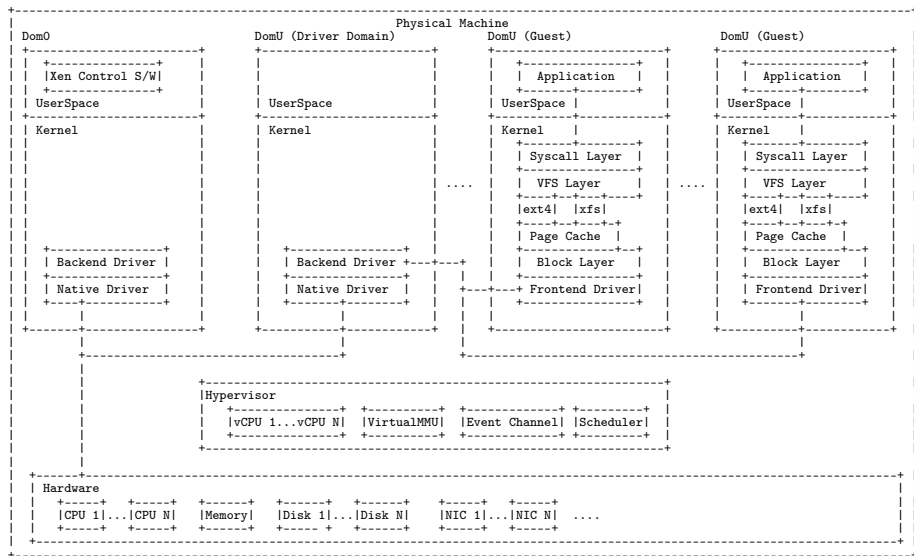


- Linux file system architecture



- Typical/Simplified workflow of how disk writes are done (assuming without `O_DIRECT` or `O_SYNC`)
  - Application makes `write(somedata)` syscall
  - Data is transferred to kernel page cache (page cache == Unused RAM is used to cache the data read/written)
  - `write()` call returns
  - After 'sometime', data is transferred to disk
  - Notes
    - The `write()` call can block at times. For example, when the page cache is full (`vm.dirty_ratio` in below example)
    - The data is written asynchronously. How often is based on some parameters and/or cache/buffer status.
    - Either writeback time is reached (`vm.dirty_writeback_centisecs` or `vm.dirtytime_expire_seconds`) or the page cache is full.
- Linux file system architecture when virtualized (xen)

# Resource: Disk III



- Impact of virtualization <sup>6</sup>

- Virtualization adds extra layer of redirection, increasing latency and bottleneck
- Additionally, network based disks (AWS EBS, EFS etc) bring-in variation in performance that we cannot control and/or measure at times.
  - AWS specific: Use EBS optimized instances, supposed to have separate/dedicated NIC for EBS traffic (Ref?)
- Performance may vary significantly due to multi-tenancy / noisy-neighbor
- For example, a high IOWait time may / may not have any relation with our IOPS (== noisy-neighbor saturating local disk controller or network card in case of network file system)

- Life of a byte in Disk IO

We will write one byte into a file and then follow that byte as it flows through various subsystems

# Resource: Disk V

```
# Life of a byte

# Install kernel debug info

yum-config-manager --enable "amzn-main-debuginfo" --enable "amzn-updates-debuginfo"
yum -y install kernel-debuginfo kernel-devel

# Setup the device
mkfs.ext4 /dev/xvdc
mount /dev/xvdc /tmp/test

# Check the block size
tune2fs -l /dev/xvdc | grep -i 'block size'
Block size:          4096

# Create a file with just one byte
echo -n "n" > /tmp/test/foo

#####

# Create probe points: vfs layer, block layer and then from xen-blkfront driver ("drivers/block/xen-blkfront.c")
perf probe --add='vfs_*' --add='blkif_*' --add='blkfront_*' --add='blkback_*' --add='xlxbd_*'

# Run the file write operation with probes enabled
# We will open the above file with O_SYNC flag, and then just update one byte in it
perf trace -T --event 'block:*' --event='probe:vfs_*' --event='probe:blkif_*' --event='probe:blkfront_*' \
--event='probe:blkback_*' --event='probe:xlxbd_*' --event='ext4:*' \
  ruby -e 'f=open("/tmp/test/foo", File::RDWR + File::SYNC); f.write("y"); f.close()'

78224651.658 ( 0.031 ms): ruby/5215 brk(                                ) = 0x11e8000
78224651.719 ( 0.033 ms): ruby/5215 mmap(len: 4096, prot: READ|WRITE, flags: PRIVATE|ANONYMOUS, fd: -1 ) = 0x7fd9b3080000
78224651.774 ( 0.020 ms): ruby/5215 access(filename: 0xb2e7e490, mode: R ) = -1 ENOENT No such file
78224651.814 ( 0.020 ms): ruby/5215 open(filename: 0xb2e7cd95, flags: CLOEXEC ) ...
.....
.....
78224789.337 ( 0.031 ms): ruby/5215 open(filename: 0x15e2810, flags: CLOEXEC|RDWR|SYNC|0x101000 ) ...
78224789.368 (           ): probe:vfs_open:(ffffff811f87b0))
78224789.337 ( 0.062 ms): ruby/5215 ... [continued]: open() = 7
78224789.426 ( 0.025 ms): ruby/5215 fcntl(fd: 7, cmd: GETFD, arg: 7 ) = 1
78224789.481 ( 0.026 ms): ruby/5215 fstat(fd: 7, statbuf: 0x7ffccbc42f0 ) ...
78224789.508 (           ): probe:vfs_fstat:(ffffff811fefc0))
78224789.532 (           ): probe:vfs_getattr:(ffffff811fef90))
```

# Resource: Disk VI

```
78224789.558 (      ): probe:vfs_getattr_nosec:(ffffff811fee60))
78224789.481 ( 0.108 ms): ruby/5215 ... [continued]: fstat() = 0
78224789.614 ( 0.025 ms): ruby/5215 ioctl(fd: 7, cmd: TCGETS, arg: 0x7ffccbc4340 ) = -1 ENOTTY Inappropriate
78224789.670 ( 0.029 ms): ruby/5215 write(fd: 7, buf: 0x15e3fe0, count: 1 ) ...
78224789.699 (      ): probe:vfs_write:(ffffff811fa070))
78224789.731 (      ): ext4:ext4_journal_start:dev 202,32 blocks, 2 rsv_blocks, 0 caller ext4_dirty_inode)
78224789.762 (      ): ext4:ext4_mark_inode_dirty:dev 202,32 ino 12 caller ext4_dirty_inode)
78224789.791 (      ): block:block_touch_buffer:202,32 sector=1057 size=4096)
78224789.827 (      ): ext4:ext4_da_write_begin:dev 202,32 ino 12 pos 0 len 1 flags 0)
78224789.855 (      ): ext4:ext4_journal_start:dev 202,32 blocks, 1 rsv_blocks, 0 caller ext4_da_write_begin)
78224789.885 (      ): ext4:ext4_da_write_end:dev 202,32 ino 12 pos 0 len 1 copied 1)
78224789.914 (      ): block:block_dirty_buffer:202,32 sector=34304 size=4096)
78224789.967 (      ): probe:vfs_fsync_range:(ffffff8122c1f0))
78224789.990 (      ): ext4:ext4_sync_file_enter:dev 202,32 ino 12 parent 2 datasync 0 )
78224790.020 (      ): ext4:ext4_writepages:dev 202,32 ino 12 nr_to_write 9223372036854775807 pages_skipped 0 range_start 0 range_
78224790.048 (      ): ext4:ext4_journal_start:dev 202,32 blocks, 8 rsv_blocks, 0 caller ext4_writepages)
78224790.109 (      ): ext4:ext4_da_write_pages:dev 202,32 ino 12 first_page 0 nr_to_write 9223372036854775807 sync_mode 1)
78224790.142 (      ): block:block_bio_queue:202,32 WS 274432 + 8 [ruby])
78224790.170 (      ): block:block_getrq:202,32 WS 274432 + 8 [ruby])
78224790.199 (      ): block:block_plug:[ruby])
78224790.221 (      ): block:block_rq_insert:202,32 WS 0 () 274432 + 8 [ruby])
78224790.235 (      ): block:block_unplug:[ruby] 1)
78224790.237 (      ): block:block_rq_issue:202,32 WS 0 () 274432 + 8 [ruby])
78224790.238 (      ): probe:blkif_queue_request:(ffffff8143f970))
78224790.239 (      ): probe:blkif_ring_get_request:(ffffff8143d810))
78224790.240 (      ): probe:blkif_setup_rw_req_grant:(ffffff814435d0))
78224790.315 (      ): ext4:ext4_writepages_result:dev 202,32 ino 12 ret 0 pages_written 1 pages_skipped 0 sync_mode 1 writeback_i
.....
78224792.127 (      ): ext4:ext4_sync_file_exit:dev 202,32 ino 12 ret 0)
78224789.670 ( 2.503 ms): ruby/5215 ... [continued]: write() = 1
78224792.203 ( 0.030 ms): ruby/5215 close(fd: 7 ) = 0
.....
78224803.693 ( 0.000 ms): ruby/5215 exit_group( )

#####

# Looking above operations from PoV of block layer

btrace /dev/xvdc
202,32 0 1 43.474871183 5657 Q R 270344 + 8 [ruby]
```

# Resource: Disk VII

```
202,32 0 2 43.474872807 5657 G R 270344 + 8 [ruby]
202,32 0 3 43.474873303 5657 I R 270344 + 8 [ruby]
202,32 0 4 43.474873914 5657 D R 270344 + 8 [ruby]
202,32 2 2 43.475358283 0 C R 270344 + 8 [0]
202,32 2 3 43.475382340 5657 Q WS 270344 + 8 [ruby]
202,32 2 4 43.475383344 5657 G WS 270344 + 8 [ruby]
202,32 2 5 43.475383618 5657 P N [ruby]
202,32 2 6 43.475384342 5657 I WS 270344 + 8 [ruby]
202,32 2 7 43.475384678 5657 U N [ruby] 1
202,32 2 8 43.475384971 5657 D WS 270344 + 8 [ruby]
202,32 2 9 43.475922196 0 C WS 270344 + 8 [0]
202,32 2 10 43.475931892 5657 Q WSM 8456 + 8 [ruby]
202,32 2 11 43.475932544 5657 G WSM 8456 + 8 [ruby]
202,32 2 12 43.475932855 5657 I WSM 8456 + 8 [ruby]
202,32 2 13 43.475933142 5657 D WSM 8456 + 8 [ruby]
202,32 2 14 43.476441001 0 C WSM 8456 + 8 [0]
```

So looks like 8 sectors, starting from 270344 were written.  
Why?

# Lets get the file details

```
stat /tmp/test/foo
  File: '/tmp/test/foo'
  Size: 1          Blocks: 8          IO Block: 4096   regular file
Device: ca20h/51744d Inode: 12       Links: 1
.....
.....
```

# So it is 8 sectors (because block size is 4096)

# Get the file's sector details

hdparm --fibmap /tmp/test/foo

```
/tmp/test/foo:
filesystem blocksize 4096, begins at LBA 0; assuming 512 byte sectors.
byte_offset begin_LBA end_LBA sectors
0 270344 270351 8
```

# So it is indeed sector 270344 that is where the file is stored

# Check if we have the single byte 'y' we wrote stored there.

dd if=/dev/xvdc bs=512 skip=270344 count=1 status=none | hexdump -C

# Resource: Disk VIII

```
00000000 79 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |y.....|
00000010 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 |.....|
*
00000200
```

## • Disk utilization at system level / per device

Check system load, high load avg might indicate disk utilization/saturation as well

# uptime

04:48:56 up 173 days, 21:51, 2 users, load average: 85.01, 84.57, 83.81

#####

Check disk utilization/saturation by device

# iostat -xz 1

Linux 4.4.44-39.55.amzn1.x86\_64 (cluster-2-data-108) 08/11/2017 \_x86\_64\_ (8 CPU)

avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
	1.56	0.02	0.18	0.66	0.01	97.57

Device:	rrqm/s	wrqm/s	r/s	w/s	rsec/s	wsec/s	avgrq-sz	avgqu-sz	await	svctm	%util
xvda	0.00	2.68	0.07	2.31	2.59	42.58	18.91	0.00	1.65	0.95	0.23
xvdi	1.47	1.07	13.86	16.79	1121.99	2295.66	111.52	0.19	6.18	0.97	2.97
xvdj	1.47	1.31	13.86	18.66	1121.92	2311.07	105.56	0.02	6.64	0.96	3.11
md0	0.00	0.00	26.90	53.28	2243.91	4606.73	85.44	0.00	0.00	0.00	0.00
dm-0	0.00	0.00	20.97	36.57	2243.91	4606.73	119.07	0.08	3.67	1.01	5.79

#####



# Resource: Disk IX

```
# dstat 1
----total-cpu-usage---- -dsk/total- -net/total- ---paging-- ---system--
usr sys idl wai hiq siq| read  writ| recv  send|  in   out | int   csw
 2   0   98   1   0   0|1125k 2346k|    0    0|    0    0 | 980 1253
 2   0   0   98   0   0|    0  160k|6642B 4859B|    0    0 | 881 1173
 0   0   0  100   0   0|    0    0|   66B  158B|    0    0 | 245  448
 0   0   0  100   0   0|    0    0|  426B 6394B|    0    0 | 295  510
 0   0   0  100   0   0|    0    0|   164B 216B|    0    0 | 252  457
 0   0   0  100   0   0|    0    0|   737B 632B|    0    0 | 698  906
 0   0   0  100   0   0|    0    0|   164B 216B|    0    0 | 594  888
 0   0   0  100   0   0|    0    0|2355B 2426B|    0    0 | 411  633
 0   0   0  100   0   0|    0    0|   328B 7872B|    0    0 | 303  511
 0   0   0  100   0   0|    0    0|    66B  126B|    0    0 | 710  939
 1   1   0   98   0   0|    0  208k|2643B 1566B|    0    0 |2063 2309
 1   0   0   99   0   0|    0  2136B 2313B|    0    0 | 810 1145
 0   0   0  100   0   0|    0    0|   295B 1428B|    0    0 | 801  955
 0   0   0  100   0   0|    0    0|   240B 7806B|    0    0 | 292  508
 0   0   0  100   0   0|    0    0|2266B 2297B|    0    0 | 402  638
 0   0   0  100   0   0|    0  88k| 639B  534B|    0    0 | 560  830
 0   0   0  100   0   0|    0 280k|   66B  134B|    0    0 | 814 1066
```

#####

```
# vmstat 1
procs -----memory----- ---swap-- ----io---- --system-- ----cpu-----
 r  b   swpd   free   buff  cache   si   so    bi   bo   in   cs us sy id wa st
0 21      0 452900  99716 27156232    0    0   141   291    0    0  2  0 98  1  0
0 21      0 453132  99716 27156232    0    0    0    0 452  743  0  0  0 100  0
0 21      0 453132  99720 27156232    0    0    0   48 267  480  0  0  0 100  0
```

# Resource: Disk X

0	21	0	452460	99720	27156232	0	0	0	0	1261	1403	0	0	0	99	0
0	21	0	452512	99720	27156232	0	0	0	0	715	944	0	0	0	100	0
0	21	0	452512	99720	27156232	0	0	0	4	320	535	0	0	0	100	0
0	21	0	452636	99720	27156232	0	0	0	0	504	834	0	0	0	100	0
0	21	0	452636	99724	27156228	0	0	0	20	251	457	0	0	0	100	0
0	21	0	452636	99724	27156232	0	0	0	0	321	484	0	0	0	100	0
0	21	0	452636	99724	27156232	0	0	0	0	303	529	0	0	0	100	0
0	21	0	452760	99724	27156232	0	0	0	0	705	941	0	0	0	100	0
0	21	0	452124	99724	27156232	0	0	0	0	1192	1463	0	0	0	99	0
0	21	0	452124	99724	27156232	0	0	0	60	265	472	0	0	0	100	0
0	21	0	452008	99724	27156232	0	0	0	16	920	1051	0	1	0	99	0
0	21	0	452008	99724	27156232	0	0	0	0	348	603	0	0	0	100	0
0	21	0	452008	99724	27156232	0	0	0	0	242	442	0	0	0	100	0

- Disk utilization by process

# Resource: Disk XI

```
# iotop -o
```

```
Total DISK READ: 0.00 B/s | Total DISK WRITE: 7.94 K/s
```

TID	PRI	USER	DISK READ	DISK WRITE	SWAPIN	IO>	COMMAND
7720	be/4	deploy	0.00 B/s	3.97 K/s	0.00 %	0.00 %	Passenger RubyApp: /data/helpkit/current/public (pr
25654	be/4	root	0.00 B/s	47.65 K/s	0.00 %	0.00 %	Passenger core
25655	be/4	root	0.00 B/s	0.00 B/s	0.00 %	0.00 %	Passenger core
21647	be/4	root	0.00 B/s	3.97 K/s	0.00 %	0.00 %	Passenger core
21229	be/4	deploy	0.00 B/s	3.97 K/s	0.00 %	0.00 %	nginx: worker process
20932	be/4	deploy	0.00 B/s	35.74 K/s	0.00 %	0.00 %	Passenger RubyApp: /data/helpkit/current/public (pr
15151	be/4	deploy	0.00 B/s	7.94 K/s	0.00 %	0.00 %	Passenger RubyApp: /data/helpkit/current/public (pr
25656	be/4	root	0.00 B/s	55.59 K/s	0.00 %	0.00 %	Passenger core
25657	be/4	root	0.00 B/s	51.62 K/s	0.00 %	0.00 %	Passenger core
23081	be/4	deploy	0.00 B/s	51.62 K/s	0.00 %	0.00 %	Passenger RubyApp: /data/helpkit/current/public (pr
16853	be/4	root	0.00 B/s	7.94 K/s	0.00 %	0.00 %	java -XX:+UseParallelGC -server -Xms64m -Xmx128m -D

```
#####
```

```
# pidstat -d
```

```
Linux 4.4.51-40.67.amzn1.x86_64 (rails-app-4) 08/09/2017 _x86_64_ (4 CPU)
```

07:54:05 AM	PID	kB_rd/s	kB_wr/s	kB_ccwr/s	Command
07:54:05 AM	1	36.06	39.42	6.94	init
07:54:05 AM	31	0.00	0.00	0.00	xenwatch
07:54:05 AM	1571	0.00	3.95	0.00	jbd2/xvda1-8
07:54:05 AM	1614	0.00	0.00	0.00	udev
07:54:05 AM	2439	0.00	0.00	0.00	dhclient
07:54:05 AM	2548	0.00	0.00	0.00	dhclient
07:54:05 AM	2595	0.00	0.20	0.00	auditd
.....					
07:53:36 AM	21204	0.00	0.00	0.00	PassengerAgent

# Resource: Disk XII

07:53:36 AM	21207	0.00	29.44	28.89	PassengerAgent
07:53:36 AM	21214	0.00	0.00	0.00	PassengerAgent
07:53:36 AM	21223	0.00	0.00	0.00	nginx
07:53:36 AM	21230	0.00	0.42	0.28	nginx
07:53:36 AM	21231	0.00	0.41	0.26	nginx
07:53:36 AM	21234	0.00	0.44	0.29	nginx
07:53:36 AM	21235	0.00	0.43	0.28	nginx
07:53:36 AM	21236	0.00	0.38	0.23	nginx
07:53:36 AM	21253	0.50	37.51	2.33	ruby
07:53:36 AM	23074	0.00	0.03	0.00	ruby
07:53:36 AM	25853	0.00	0.00	0.00	pidstat

- Find processes that are in uninterruptible state (most likely due to disk IO)

# Resource: Disk XIII

Processes that are in uninterruptible state

```
# ps axl | awk ' $10 ~ /[D]/ {
0 0 654 653 20 0 118448 1548 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
1 0 677 2 20 0 0 0 - D ? 59:19 [kswapd0]
0 0 1024 1021 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 1255 1251 20 0 118448 1428 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
1 0 1319 2 20 0 0 0 - D ? 0:12 [kworker/1:0]
1 0 1671 2 0 -20 0 0 - D< ? 0:23 [kworker/2:2H]
0 0 1933 1929 20 0 118448 1464 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 3458 3453 20 0 118448 1536 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 4058 4057 20 0 118448 1424 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 4159 4156 20 0 118448 1536 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 5739 5737 20 0 118448 1548 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 5775 5773 20 0 118448 1516 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 5786 5781 20 0 118448 1460 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 6293 6292 20 0 118448 1520 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 6501 6497 20 0 118448 1452 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
1 0 6686 2 20 0 0 0 - D ? 47:52 [xfsaild/dm-0]
0 0 8147 8142 20 0 118448 1424 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 8589 8586 20 0 118448 1548 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 8784 8779 20 0 118448 1536 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 9463 9460 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 10988 10986 20 0 118448 1460 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 11606 11603 20 0 118448 1584 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 11695 11694 20 0 118448 1520 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 12189 12188 20 0 118448 1584 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 12288 12283 20 0 118448 1584 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 13248 13246 20 0 118448 1552 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 13932 13929 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 14153 14151 20 0 118448 1424 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
1 0 15115 2 20 0 0 0 - D ? 0:09 [kworker/2:2]
0 0 15724 15721 20 0 118448 1532 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 16157 16156 20 0 118448 1584 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 16374 16371 20 0 118448 1456 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 17458 17453 20 0 118448 1512 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 18530 18529 20 0 118448 1524 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 18591 18588 20 0 118448 1552 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 18628 18625 20 0 118448 1460 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 18637 18632 20 0 118448 1372 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 19324 19319 20 0 118448 1460 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 20835 20833 30 10 120572 2316 - DN ? 0:00 /usr/sbin/logrotate /etc/logrotate.conf
```

# Resource: Disk XIV

```
0 0 20877 20876 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 21452 21450 20 0 118448 1528 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 21624 21619 20 0 118448 1380 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 23113 23110 20 0 118448 1520 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 23676 23673 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 23881 23876 20 0 118448 1516 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 25271 25270 20 0 118448 1368 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 25477 25474 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 25828 25826 20 0 118448 1512 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 26043 26041 20 0 118448 1548 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 26167 26162 20 0 118448 1380 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 26774 26773 20 0 118448 1512 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 28326 28325 20 0 118448 1520 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 28981 28978 20 0 118448 1480 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 28982 28977 20 0 118448 1584 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 30585 30584 20 0 118448 1516 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 30912 30909 20 0 118448 2196 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 31209 31206 20 0 118448 1512 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 31456 31451 20 0 118448 1532 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
0 0 31816 31812 20 0 118448 1588 - Ds ? 0:00 /usr/sbin/logrotate /etc/logrotate.d/goaudit
```

#####

What are they waiting on?

```
# ps axl | awk '$10 ~ /[D]/' | sudo awk '{ print "=====$13=====$3===="; system("cat /proc/"$3"/stack")}'
====/usr/sbin/logrotate===654=====
[<ffffffff812daa54>] call_rwsem_down_read_failed+0x14/0x30
[<ffffffffa03b303f>] xfs_ilock+0xff/0x130 [xfs]
[<ffffffffa03b30a0>] xfs_ilock_data_map_shared+0x30/0x40 [xfs]
[<ffffffffa03a7060>] xfs_dir_open+0x30/0x60 [xfs]
[<ffffffff811d2e23>] do_dentry_open+0x223/0x300
[<ffffffff811d40e5>] vfs_open+0x55/0x80
[<ffffffff811e17e0>] path_openat+0x1b0/0x12a0
[<ffffffff811e467e>] do_filp_open+0x7e/0xd0
[<ffffffff811d4468>] do_sys_open+0x128/0x210
[<ffffffff811d4584>] SyS_openat+0x14/0x20
[<ffffffff814efcae>] entry_SYSCALL_64_fastpath+0x12/0x71
[<ffffffffffffffff>] 0xffffffffffffffff
====[kswapd0]===677=====
[<ffffffff81083f5f>] flush_work+0xef/0x170
[<ffffffffa03c5659>] xlog_cil_force_lsn+0x79/0x1e0 [xfs]
```

# Resource: Disk XV

```
[<fffffffa03c3cd1>] _xfs_log_force_lsn+0x71/0x310 [xfs]
[<fffffffa03c3f9e>] xfs_log_force_lsn+0x2e/0xa0 [xfs]
[<fffffffa03b2b8d>] __xfs_iunpin_wait+0x8d/0x140 [xfs]
[<fffffffa03b6329>] xfs_iunpin_wait+0x19/0x20 [xfs]
[<fffffffa03ab722>] xfs_reclaim_inode+0x122/0x340 [xfs]
[<fffffffa03abb54>] xfs_reclaim_inodes_ag+0x214/0x330 [xfs]
[<fffffffa03ac773>] xfs_reclaim_inodes_nr+0x33/0x40 [xfs]
[<fffffffa03bb099>] xfs_fs_free_cached_objects+0x19/0x20 [xfs]
[<ffffff811d85c1>] super_cache_scan+0x181/0x190
[<ffffff81172a56>] shrink_slab.part.41+0x206/0x3f0
[<ffffff81176809>] shrink_zone+0x2a9/0x2c0
[<ffffff81177794>] kswapd+0x4b4/0x960
[<ffffff8108a7a9>] kthread+0xc9/0xe0
[<ffffff814f000f>] ret_from_fork+0x3f/0x70
[<fffffffffffffff>] 0xfffffffffffffff
=====usr/sbin/logrotate===1024=====
[<ffffff812daa54>] call_rwsem_down_read_failed+0x14/0x30
[<fffffffa03b303f>] xfs_ilock+0xff/0x130 [xfs]
[<fffffffa03b30a0>] xfs_ilock_data_map_shared+0x30/0x40 [xfs]
[<fffffffa03a7060>] xfs_dir_open+0x30/0x60 [xfs]
[<ffffff811d2e23>] do_dentry_open+0x223/0x300
[<ffffff811d40e5>] vfs_open+0x55/0x80
[<ffffff811e17e0>] path_openat+0x1b0/0x12a0
[<ffffff811e467e>] do_filp_open+0x7e/0xd0
[<ffffff811d4468>] do_sys_open+0x128/0x210
[<ffffff811d4584>] SyS_openat+0x14/0x20
[<ffffff814efcae>] entry_SYSCALL_64_fastpath+0x12/0x71
[<fffffffffffffff>] 0xfffffffffffffff
=====usr/sbin/logrotate===1255=====
[<ffffff812daa54>] call_rwsem_down_read_failed+0x14/0x30
[<fffffffa03b303f>] xfs_ilock+0xff/0x130 [xfs]
[<fffffffa03b30a0>] xfs_ilock_data_map_shared+0x30/0x40 [xfs]
[<fffffffa03a7060>] xfs_dir_open+0x30/0x60 [xfs]
[<ffffff811d2e23>] do_dentry_open+0x223/0x300
[<ffffff811d40e5>] vfs_open+0x55/0x80
[<ffffff811e17e0>] path_openat+0x1b0/0x12a0
[<ffffff811e467e>] do_filp_open+0x7e/0xd0
[<ffffff811d4468>] do_sys_open+0x128/0x210
[<ffffff811d4584>] SyS_openat+0x14/0x20
[<ffffff814efcae>] entry_SYSCALL_64_fastpath+0x12/0x71
[<fffffffffffffff>] 0xfffffffffffffff
=====kworker/1:0===1319=====
```

# Resource: Disk XVI

```
[<ffffff810b2b51>] down+0x41/0x50
[<ffffffa03a2afc>] xfs_buf_lock+0x3c/0xf0 [xfs]
[<ffffffa03a2d12>] _xfs_buf_find+0x162/0x340 [xfs]
[<ffffffa03a2f1a>] xfs_buf_get_map+0x2a/0x280 [xfs]
[<ffffffa03a3bdd>] xfs_buf_read_map+0x2d/0x180 [xfs]
[<ffffffa03cf664>] xfs_trans_read_buf_map+0xf4/0x310 [xfs]
[<ffffffa037c329>] xfs_btree_read_buf_block.constprop.28+0x69/0xa0 [xfs]
[<ffffffa037c3d1>] xfs_btree_lookup_get_block+0x71/0xe0 [xfs]
[<ffffffa0380c37>] xfs_btree_lookup+0xb7/0x560 [xfs]
[<ffffffa0367091>] xfs_free_ag_extent+0x61/0x760 [xfs]
[<ffffffa03687ea>] xfs_free_extent+0xda/0x110 [xfs]
[<ffffffa03cfff16>] xfs_trans_free_extent+0x26/0x60 [xfs]
[<ffffffa039f74f>] xfs_bmap_finish+0xff/0x120 [xfs]
[<ffffffa03b5453>] xfs_itruncate_extents+0x113/0x240 [xfs]
[<ffffffa03a0384>] xfs_free_eofblocks+0x1b4/0x210 [xfs]
[<ffffffa03accf5>] xfs_inode_free_eofblocks+0x95/0x160 [xfs]
[<ffffffa03ab2ce>] xfs_inode_ag_walk.isra.10+0x1ee/0x310 [xfs]
[<ffffffa03ac531>] xfs_inode_ag_iterator_tag+0x71/0xa0 [xfs]
[<ffffffa03ac7fd>] xfs_icache_free_eofblocks+0x2d/0x40 [xfs]
[<ffffffa03ac82b>] xfs_eofblocks_worker+0x1b/0x30 [xfs]
[<ffffff81084ba0>] process_one_work+0x150/0x3f0
[<ffffff8108531a>] worker_thread+0x11a/0x470
[<ffffff8108a7a9>] kthread+0xc9/0xe0
[<ffffff814f000f>] ret_from_fork+0x3f/0x70
[<fffffffffffff>] 0xfffffffffffffff
====[kworker/2:2H]===1671====
[<ffffff81083f5f>] flush_work+0xef/0x170
[<ffffffa03c5659>] xlog_cil_force_lsn+0x79/0x1e0 [xfs]
[<ffffffa03c3986>] _xfs_log_force+0x76/0x270 [xfs]
[<ffffffa03c3ba6>] xfs_log_force+0x26/0x90 [xfs]
[<ffffffa03c3c34>] xfs_log_worker+0x24/0x50 [xfs]
[<ffffff81084ba0>] process_one_work+0x150/0x3f0
[<ffffff8108531a>] worker_thread+0x11a/0x470
[<ffffff8108a7a9>] kthread+0xc9/0xe0
[<ffffff814f000f>] ret_from_fork+0x3f/0x70
[<fffffffffffff>] 0xfffffffffffffff
.....
.....
=====/usr/sbin/logrotate===31816====
[<ffffff812daa54>] call_rwsem_down_read_failed+0x14/0x30
[<ffffffa03b303f>] xfs_ilock+0xff/0x130 [xfs]
```



# Resource: Disk XVII

```
<fffffffa03b30a0> xfs_iloc_data_map_shared+0x30/0x40 [xfs]
<fffffffa03a7060> xfs_dir_open+0x30/0x60 [xfs]
<ffffff811d2e23> do_dentry_open+0x223/0x300
<ffffff811d40e5> vfs_open+0x55/0x80
<ffffff811e17e0> path_openat+0x1b0/0x12a0
<ffffff811e467e> do_filp_open+0x7e/0xd0
<ffffff811d4468> do_sys_open+0x128/0x210
<ffffff811d4584> SyS_openat+0x14/0x20
<ffffff814efcae> entry_SYSCALL_64_fastpath+0x12/0x71
<ffffffffffffff> 0xffffffffffffff
```

## • Find disk activity at block IO layer

Find the activity at block layer level

NOTE: Use btt tool for extended analysis: <http://www.cse.unsw.edu.au/~aaronc/iosched/doc/btt.html>

```
# btrace == blktrace /dev/xvda -o - | blkparse -s -i -
# btrace /dev/xvda
202,0 3 1 0.000000000 11720 A W 6257952 + 8 <- (202,1) 6253856
202,0 3 2 0.000000904 11720 Q W 6257952 + 8 [java]
202,0 3 3 0.000004811 11720 G W 6257952 + 8 [java]
202,0 3 4 0.000005466 11720 P N [java]
.....
202,0 3 55 1.337693195 1571 A WS 4584696 + 8 <- (202,1) 4580600
202,0 3 56 1.337693283 1571 Q WS 4584696 + 8 [jbd2/xvda1-8]
202,0 3 57 1.337693403 1571 M WS 4584696 + 8 [jbd2/xvda1-8]
202,0 3 58 1.337693622 1571 A WS 4584704 + 8 <- (202,1) 4580608
202,0 3 59 1.337693710 1571 Q WS 4584704 + 8 [jbd2/xvda1-8]
.....
202,0 3 126 2.001851710 11720 Q W 6256664 + 8 [java]
202,0 3 127 2.001855937 11720 G W 6256664 + 8 [java]
202,0 3 128 2.001856577 11720 P N [java]
202,0 3 129 2.001859526 11720 I W 6256664 + 8 [java]
202,0 3 130 2.001860479 11720 U N [java] 1
202,0 3 131 2.001861857 11720 D W 6256664 + 8 [java]
```

# Resource: Disk XVIII

```
202,0 3 132 2.002029389 11720 A W 6261568 + 8 <- (202,1) 6257472
202,0 3 133 2.002029792 11720 Q W 6261568 + 8 [java]
202,0 3 134 2.002031071 11720 G W 6261568 + 8 [java]
202,0 3 135 2.002031370 11720 P N [java]
202,0 3 136 2.002032645 11720 I W 6261568 + 8 [java]
202,0 3 137 2.002033044 11720 U N [java] 1
202,0 3 138 2.002033469 11720 D W 6261568 + 8 [java]
202,0 3 139 2.002453955 0 C W 6256664 + 8 [0]
202,0 3 140 2.002516093 0 C W 6261568 + 8 [0]
202,0 3 141 3.002859806 11720 A W 6255424 + 8 <- (202,1) 6251328
202,0 3 142 3.002860373 11720 Q W 6255424 + 8 [java]
202,0 3 143 3.002862233 11720 G W 6255424 + 8 [java]
202,0 3 144 3.002862542 11720 P N [java]
202,0 3 145 3.002864546 11720 I W 6255424 + 8 [java]
202,0 3 146 3.002864960 11720 U N [java] 1
```

.....

```
202,0 3 147 3.002865474 11720 D W 6255424 + 8 [java]
202,0 3 268 10.009390888 11720 D W 6257960 + 8 [java]
202,0 3 269 10.009769306 0 C W 6257952 + 8 [0]
202,0 3 270 10.009931136 0 C W 6257960 + 8 [0]
```

~C

```
java (11720)
Reads Queued: 0, 0KiB Writes Queued: 22, 88KiB
Read Dispatches: 0, 0KiB Write Dispatches: 22, 88KiB
Reads Requeued: 0, 0KiB Writes Requeued: 0, 0KiB
Reads Completed: 0, 0KiB Writes Completed: 0, 0KiB
Read Merges: 0, 0KiB Write Merges: 0, 0KiB
IO unplugs: 22, 0KiB Timer unplugs: 0, 0KiB
Allocation wait: 0, 0KiB Allocation wait: 0, 0KiB
Dispatch wait: 0, 0KiB Dispatch wait: 0, 0KiB
Completion wait: 0, 0KiB Completion wait: 0, 0KiB
```

```
jbd2/xvda1-8 (1571)
Reads Queued: 0, 0KiB Writes Queued: 46, 184KiB
Read Dispatches: 0, 0KiB Write Dispatches: 4, 184KiB
Reads Requeued: 0, 0KiB Writes Requeued: 0, 0KiB
Reads Completed: 0, 0KiB Writes Completed: 0, 0KiB
Read Merges: 0, 0KiB Write Merges: 42, 168KiB
IO unplugs: 2, 0KiB Timer unplugs: 0, 0KiB
Allocation wait: 0, 0KiB Allocation wait: 0, 0KiB
Dispatch wait: 0, 0KiB Dispatch wait: 0, 0KiB
Completion wait: 0, 0KiB Completion wait: 0, 0KiB
```

# Resource: Disk XIX

```
swapper/3 (0)
Reads Queued:      0,      OKiB Writes Queued:      0,      OKiB
Read Dispatches:   0,      OKiB Write Dispatches:   0,      OKiB
Reads Requeued:    0,      Writes Requeued:    0
Reads Completed:   0,      OKiB Writes Completed:   24,      264KiB
Read Merges:       0,      OKiB Write Merges:       0,      OKiB
IO unplugs:        0,      Timer unplugs:      0
Allocation wait:    0,      Allocation wait:     0
Dispatch wait:      0,      Dispatch wait:       0
Completion wait:    0,      Completion wait:     0

utils.rb:110 (6680)
Reads Queued:      0,      OKiB Writes Queued:      0,      OKiB
Read Dispatches:   0,      OKiB Write Dispatches:   0,      OKiB
Reads Requeued:    0,      Writes Requeued:    0
Reads Completed:   0,      OKiB Writes Completed:   2,      8KiB
Read Merges:       0,      OKiB Write Merges:       0,      OKiB
IO unplugs:        0,      Timer unplugs:      0
Allocation wait:    0,      Allocation wait:     0
Dispatch wait:      0,      Dispatch wait:       0
Completion wait:    0,      Completion wait:     0

.....
.....

Throughput (R/W): 0KiB/s / 27KiB/s
Events (202,0): 330 entries
Skips: 0 forward (0 - 0.0%)
```

- Using blktrace to trace/observe the activity at block layer

# Resource: Disk XX

```
$ btrace /dev/xvdz
202,6400 1 1 0.0000 15982 Q R 0 + 32 [dd]
202,6400 1 2 0.0000 15982 G R 0 + 32 [dd]
202,6400 1 3 0.0000 15982 P N [dd]
202,6400 1 4 0.0000 15982 I R 0 + 32 [dd]
202,6400 1 5 0.0000 15982 U N [dd] 1
202,6400 1 6 0.0000 15982 D R 0 + 32 [dd]
202,6400 3 1 0.0005 0 C R 0 + 32 [0]

|| $ echo 3 > /proc/sys/vm/drop_caches
|| $ dd if=/dev/xvdz bs=512 of=/dev/null count=1
|| 1+0 records in
|| 1+0 records out
|| 512 bytes (512 B) copied, 0.000728468 s, 703 kB/s
||
|| $ dd if=/dev/xvdz bs=512 of=/dev/null count=2
|| 2+0 records in
|| 2+0 records out
|| 1024 bytes (1.0 kB) copied, 8.07e-05 s, 12.7 MB/s
||
|| $ dd if=/dev/xvdz bs=512 of=/dev/null count=8
|| 8+0 records in
|| 8+0 records out
|| 4096 bytes (4.1 kB) copied, 0.000111268 s, 36.8 MB/s
||
||
|| $ dd if=/dev/xvdz bs=512 of=/dev/null count=9
|| 9+0 records in
|| 9+0 records out
|| 4608 bytes (4.6 kB) copied, 0.000149306 s, 30.9 MB/s
||
||
|| $ dd if=/dev/xvdz bs=512 of=/dev/null count=1 iflag=direct
|| 1+0 records in
|| 1+0 records out
|| 512 bytes (512 B) copied, 0.000728468 s, 703 kB/s
||

202,6400 0 1 32.8711 16110 Q R 32 + 64 [dd]
202,6400 0 2 32.8711 16110 G R 32 + 64 [dd]
202,6400 0 3 32.8711 16110 P N [dd]
202,6400 0 4 32.8711 16110 I R 32 + 64 [dd]
202,6400 0 5 32.8711 16110 U N [dd] 1
202,6400 0 6 32.8711 16110 D R 32 + 64 [dd]
202,6400 3 2 32.8719 0 C R 32 + 64 [0]

202,6400 2 1 147.7486 17283 Q R 0 + 1 [dd]
202,6400 2 2 147.7486 17283 G R 0 + 1 [dd]
202,6400 2 3 147.7486 17283 P N [dd]
202,6400 2 4 147.7486 17283 I R 0 + 1 [dd]
202,6400 2 5 147.7486 17283 U N [dd] 1
202,6400 2 6 147.7486 17283 D R 0 + 1 [dd]
202,6400 3 3 147.7490 9973 C R 0 + 1 [0]
```

- Example trace out of a "bad" disk

```
# perf trace --event 'block:*' dd if=/dev/xvdd of=/dev/null bs=512 count=1 iflag=direct
0.175 ( 0.016 ms): dd/28637 brk(                                ) = 0x10be000
0.221 ( 0.018 ms): dd/28637 mmap(len: 4096, prot: READ|WRITE, flags: PRIVATE|ANONYMOUS, fd: -1) = 0x7f1f56c20000
0.252 ( 0.015 ms): dd/28637 access(filename: 0x56a1f140, mode: R) = -1 ENOENT No such file or
0.282 ( 0.016 ms): dd/28637 open(filename: 0x56a1da38, flags: CLOEXEC) = 3
0.307 ( 0.012 ms): dd/28637 fstat(fd: 3, statbuf: 0x7ffc2c5e33b0) = 0
.....
1.638 ( 0.025 ms): dd/28637 open(filename: 0x2c5e5739, flags: CREAT|TRUNC|WRONLY, mode: 438) = 3
1.700 ( 0.039 ms): dd/28637 dup2(oldfd: 3, newfd: 1) = 1
1.728 ( 0.013 ms): dd/28637 close(fd: 3) = 0
1.759 ( 0.017 ms): dd/28637 clock_gettime(which_clock: MONOTONIC, tp: 0x7ffc2c5e3b40) = 0
1.809 ( 0.033 ms): dd/28637 read(buf: 0x10c0000, count: 512) = ...
1.809 (          ) : block:block_bio_queue:202,144 R 0 + 1 [dd])
1.836 (          ) : block:block_getrq:202,144 R 0 + 1 [dd])
1.858 (          ) : block:block_plug:[dd])
1.875 (          ) : block:block_rq_insert:202,144 R 0 ( ) 0 + 1 [dd])
1.887 (          ) : block:block_unplug:[dd] 1)
```

~C

- Bad disk(s) can have cascading effect on unrelated disk activity as well

# Resource: Disk XXII

```
# Ran 'yum install' on a system that had bad disk (but rootfs disk was fine)
# Yum install got stuck after about 80% work done
# Analyzing the where it is stuck showed the below stack:
#   When it tried to allocate a page out of page cache,
#   it ran out of free pages (or reached water mark), so it tried to reclaim
#   pages, which led to the trying to sync pages belonging to 'bad' disk (xfs
#   in this case), causing it to be stuck
```

```
cat /proc/'pidof yum'/stack
[<ffffffff81083f5f>] flush_work+0xef/0x170
[<ffffffffffa03c5659>] xlog_cil_force_lsn+0x79/0x1e0 [xfs]
[<ffffffffffa03c3cd1>] _xfs_log_force_lsn+0x71/0x310 [xfs]
[<ffffffffffa03c3f9e>] xfs_log_force_lsn+0x2e/0xa0 [xfs]
[<ffffffffffa03b2b8d>] __xfs_iunpin_wait+0x8d/0x140 [xfs]
[<ffffffffffa03b6329>] xfs_iunpin_wait+0x19/0x20 [xfs]
[<ffffffffffa03ab722>] xfs_reclaim_inode+0x122/0x340 [xfs]
[<ffffffffffa03abb54>] xfs_reclaim_inodes_ag+0x214/0x330 [xfs]
[<ffffffffffa03ac773>] xfs_reclaim_inodes_nr+0x33/0x40 [xfs]
[<ffffffffffa03bb099>] xfs_fs_free_cached_objects+0x19/0x20 [xfs]
[<ffffffffff811d85c1>] super_cache_scan+0x181/0x190
[<ffffffffff81172a56>] shrink_slab.part.41+0x206/0x3f0
[<ffffffffff81176809>] shrink_zone+0x2a9/0x2c0
[<ffffffffff81176ba5>] do_try_to_free_pages+0x175/0x440
[<ffffffffff81176f25>] try_to_free_pages+0xb5/0x170
[<ffffffffff8116abaa>] __alloc_pages_nodemask+0x53a/0xa60
[<ffffffffff811aef58>] alloc_pages_current+0x88/0x120
[<ffffffffff81162294>] __page_cache_alloc+0xb4/0xc0
[<ffffffffff81162c76>] pagecache_get_page+0x56/0x1e0
[<ffffffffff81162e26>] grab_cache_page_write_begin+0x26/0x40
[<ffffffffffa0120e01>] ext4_da_write_begin+0xa1/0x330 [ext4]
[<ffffffffff81161e50>] generic_perform_write+0xc0/0x1a0
[<ffffffffff81163f48>] __generic_file_write_iter+0x188/0x1e0
[<ffffffffffa0115b76>] ext4_file_write_iter+0xf6/0x360 [ext4]
[<ffffffffff811d4c5a>] __vfs_write+0xaa/0xe0
[<ffffffffff811d5282>] vfs_write+0xa2/0x1a0
[<ffffffffff811d5f86>] Sys_write+0x46/0xa0
[<ffffffffff814efcae>] entry_SYSCALL_64_fastpath+0x12/0x71
[<ffffffffffffffff>] 0xffffffffffffffff
```

- Disk space usage

```
# df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        3.7G   64K   3.7G   1% /dev
tmpfs           3.7G    0   3.7G   0% /dev/shm
/dev/xvda1      7.8G  3.4G  4.3G  44% /
/dev/xvdh       99G   5.1G   89G   6% /data
```

- Disk related errors

```
# dmesg -T | grep "blocked for more than"
INFO: task xfsaild/dm-0:6686 blocked for more than 120 seconds.
```


```
# demsg -T | grep "I/O error"
[351410.715652] EXT4-fs warning (device xvdh): htree_dirblock_to_tree:958: inode #262145: lblock 0: comm ls: error
[397736.767853] blk_update_request: I/O error, dev xvdh, sector 73992
[397736.770649] EXT4-fs warning (device xvdh): htree_dirblock_to_tree:958: inode #2: lblock 0: comm ls: error -
[399503.066719] blk_update_request: I/O error, dev xvdh, sector 73992
```

---

<sup>6</sup><http://dtrace.org/blogs/brendan/2013/01/11/virtualization-performance-zones-kvm-xen/>

# Resource: Network I

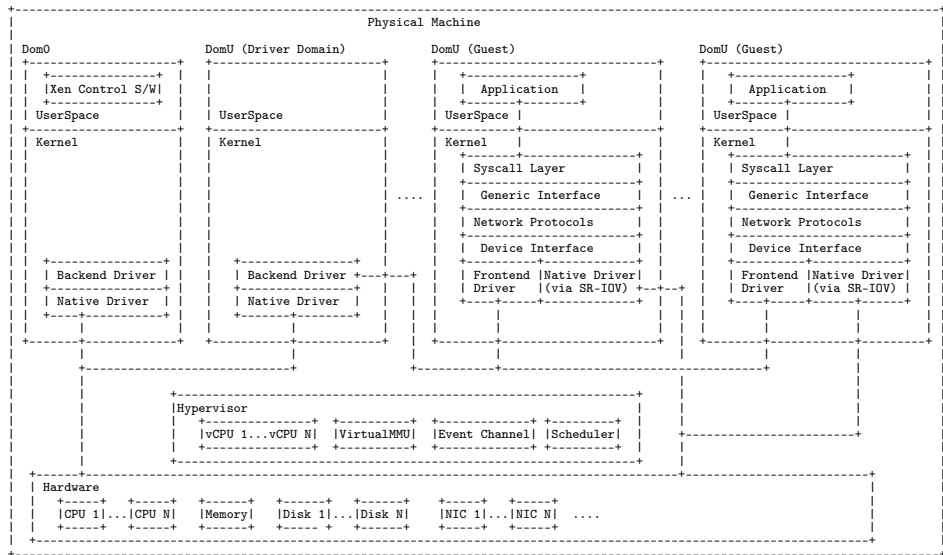
- Linux network stack architecture



- Linux when stack architecture when virtualized (xen)



# Resource: Network II



- Impact of virtualization
  - Virtualization adds extra layer of redirection, increasing latency and bottleneck
  - Performance may vary significantly due to multi-tenancy / noisy-neighbor
    - For example, sudden high latency/throughput drop may / may not have any relation with our network traffic itself (== noisy-neighbor saturating local NIC controller or switch)
  - Use SR-IOV enabled network device if available. This will allow the guest OS to directly talk to the hardware, without going through the Driver Domain
    - AWS specific: Enhanced network support

# Resource: Network IV

- Life of a byte in network stack

We will send a simple HTTP GET request and trace as it goes through various subsystems

You can find in below links more detailed walk through of various network layers: <https://blog.packagecloud.io/eng/2016/06/22/monitoring-tuning-linux-networking-stack-receiving-data/>  
<https://blog.packagecloud.io/eng/2017/02/06/monitoring-tuning-linux-networking-stack-sending-data/>

```
Install kernel debug info
# yum-config-manager --enable "amzn-main-debuginfo" --enable "amzn-updates-debuginfo"
# yum -y install kernel-debuginfo kernel-devel

Find the kernel module responsible for our ethernet card
# lspci -k
.....
00:03.0 Ethernet controller: Intel Corporation 82599 Ethernet Controller Virtual Function (rev 01)
        Kernel driver in use: ixgbevf
.....
#####

Get all the probeable functions in our driver
# perf probe -m ixgbevf -F 'ixgbevf_*'
ixgbevf_addr_list_itr
.....
```

# Resource: Network V

```
.....
ixgbevf_xmit_frame
```

Create probe point for all these functions

```
# perf probe -m ixgbevf --add='ixgbevf_*'
```

Too many( > 128) probe point found.

Added new events:

```
probe:ixgbevf_addr_list_itr (on ixgbevf_* in ixgbevf)
probe:ixgbevf_set_rx_mode (on ixgbevf_* in ixgbevf)
probe:ixgbevf_vlan_rx_kill_vid (on ixgbevf_* in ixgbevf)
probe:ixgbevf_change_mtu (on ixgbevf_* in ixgbevf)
probe:ixgbevf_set_mac (on ixgbevf_* in ixgbevf)
probe:ixgbevf_negotiate_api (on ixgbevf_* in ixgbevf)
probe:ixgbevf_free_q_vector (on ixgbevf_* in ixgbevf)
probe:ixgbevf_free_q_vectors (on ixgbevf_* in ixgbevf)
probe:ixgbevf_free_irq (on ixgbevf_* in ixgbevf)
probe:ixgbevf_update_itr (on ixgbevf_* in ixgbevf)
.....
.....
probe:ixgbevf_init_module (on ixgbevf_* in ixgbevf)
probe:ixgbevf_exit_module (on ixgbevf_* in ixgbevf)
```

You can now use it in all perf tools, such as:

```
perf record -e probe:ixgbevf_exit_module -aR sleep 1
```

```
#####
```

```
#####
```

Now trace the packets, simple ping command tracing

```
# perf trace --event 'net:*' --event 'probe:ixgbevf_*' ping -c1 8.8.8.8 >/dev/null
```

```
0.196 ( 0.027 ms): ping/3910 brk(                                ) = 0x5564c6b32000
```

```
.....
```

```
.....
```

```
5.204 ( 0.032 ms): ping/3910 sendmsg(fd: 3<socket:[87842167]>, msg: 0x5564c6a68160 ) ...
```

```
5.204 (      ) : net:net_dev_queue:dev=eth0 skbaddr=0xffff8800eaf35800 len=98)
```

```
5.229 (      ) : net:net_dev_start_xmit:dev=eth0 queue_mapping=0 skbaddr=0xffff8800eaf35800 vlan_tagged=0 vlan_proto=0x0000 vl
```

```
5.253 (      ) : probe:ixgbevf_xmit_frame:(ffffffffffa025a670))
```

```
5.276 (      ) : net:net_dev_xmit:dev=eth0 skbaddr=0xffff8800eaf35800 len=98 rc=0)
```

```
5.299 ( 0.126 ms): ping/3910 ... [continued]: sendmsg() = 64
```

# Resource: Network VI

```
5.346 ( 0.024 ms): ping/3910 setitimer(which: REAL, value: 0x7fff20936fe0) = 0
18.001 (12.632 ms): ping/3910 recvmsg(fd: 3<socket:[87842167]>, msg: 0x7fff20937000) ...
18.001 (      ) : probe:ixgbevfv_msix_clean_rings:(ffffffffffa0259b20))
18.043 (      ) : probe:ixgbevfv_poll:(ffffffffffa025c080))
.....
18.285 (12.915 ms): ping/3910 ... [continued]: recvmsg() = 84
18.294 (      ) : probe:ixgbevfv_msix_clean_rings:(ffffffffffa0259b20))
18.296 (      ) : probe:ixgbevfv_poll:(ffffffffffa025c080))
18.297 (      ) : probe:ixgbevfv_clean_rx_irq:(ffffffffffa0259b70))
18.298 (      ) : probe:ixgbevfv_update_itr:(ffffffffffa02594e0))
18.302 (      ) : net:net_dev_queue:dev=eth0 skbaddr=0xffff8801e3a0fce8 len=166)
18.304 (      ) : net:net_dev_start_xmit:dev=eth0 queue_mapping=0 skbaddr=0xffff8801e3a0fce8 vlan_tagged=0 vlan_proto=0x0000 vl
18.305 (      ) : probe:ixgbevfv_xmit_frame:(ffffffffffa025a670))
18.306 (      ) : net:net_dev_xmit:dev=eth0 skbaddr=0xffff8801e3a0fce8 len=166 rc=0)
.....
18.546 (      ) : probe:ixgbevfv_poll:(ffffffffffa025c080))
18.571 (      ) : probe:ixgbevfv_clean_rx_irq:(ffffffffffa0259b70))
18.597 (      ) : net:napi_gro_receive_entry:dev=eth0 napi_id=0x1 queue_mapping=0 skbaddr=0xffff8801e6bd4d00 vlan_tagged=0 vlan
18.621 (      ) : net:netif_receive_skb:dev=eth0 skbaddr=0xffff8801e6bd4d00 len=52)
18.649 (      ) : net:napi_gro_receive_entry:dev=eth0 napi_id=0x1 queue_mapping=0 skbaddr=0xffff8801e6bd4d00 vlan_tagged=0 vlan
18.675 (      ) : net:netif_receive_skb:dev=eth0 skbaddr=0xffff8801e6bd4d00 len=52)
18.701 (      ) : probe:ixgbevfv_alloc_rx_buffers:(ffffffffffa0259750))
.....
18.810 (      ) : probe:ixgbevfv_poll:(ffffffffffa025c080))
18.835 (      ) : probe:ixgbevfv_clean_rx_irq:(ffffffffffa0259b70))
18.860 (      ) : net:napi_gro_receive_entry:dev=eth0 napi_id=0x1 queue_mapping=0 skbaddr=0xffff8801e6bd4d00 vlan_tagged=0 vlan
18.885 (      ) : net:netif_receive_skb:dev=eth0 skbaddr=0xffff8801e6bd4d00 len=52)
.....
19.141 ( 0.024 ms): ping/3910 write(fd: 1</dev/null>, buf: 0x7fe157a07000, count: 99) = 99
19.186 ( 0.023 ms): ping/3910 write(fd: 1</dev/null>, buf: 0x7fe157a07000, count: 1) = 1
19.243 ( 0.023 ms): ping/3910 write(fd: 1</dev/null>, buf: 0x7fe157a07000, count: 145) = 145
19.265 ( 0.000 ms): ping/3910 exit_group()
```

#####

# Now we will trace a HTTP request

# We will start and establish the TCP connection, wait for keyboard input  
# and then send the HTTP GET request.

## Resource: Network VII

```
# We will only start tracing after connection establishment, so we can just
# focus on GET request alone
```

TERMINAL 1

```
(read -n 1 -p "Press any key to continue "; \
echo -e -n 'GET / HTTP/1.1\r\n' \
echo -e -n 'Host: support.freshdesk.com\r\n\r\n') \
| socat -t 10 - TCP4:support.freshdesk.com:80
Press any key to continue
```

```
HTTP/1.1 302 Found
Cache-Control: no-cache
Content-Type: text/html; charset=utf-8
Date: Tue, 05 Sep 2017 08:37:21 GMT
Location: https://support.freshdesk.com/
Set-Cookie: _x_w=1; path=/
Status: 302 Found
X-Frame-Options: SAMEORIGIN
X-Rack-Cache: miss
X-Request-Id: 794df2c16bc159a2dd339b0c33cc4394d
X-Runtime: 0.018734
X-UA-Compatible: IE=Edge,chrome=1
X-XSS-Protection: 1; mode=block
Content-Length: 96
Connection: Close
```

```
<html><body>You are being
<a href="https://support.freshdesk.com/">redirected
</a>.</body></html>
```

TERMINAL 3

```
perf trace -T --event 'net:*' --event 'probe:vfs*' --event 'probe:ixgbevf_*' -p 'pidof socat'
```

TERMINAL 2

```

||| # tshark -f "not port 22"
|||
||| Running as user "root" and group "root". This could be dangerous.
||| Capturing on eth0
|||
||| 0.000000000 172.23.3.135 -> 172.23.0.2 DNS 81 Standard query 0xc5dc A s
||| 0.000234617 172.23.0.2 -> 172.23.3.135 DNS 113 Standard query response 0
||| 0.000341919 172.23.3.135 -> 52.206.84.26 TCP 74 43074 > http [SYN] Seq=0 W
||| 0.001194278 52.206.84.26 -> 172.23.3.135 TCP 74 http > 43074 [SYN, ACK] Se
||| 0.001208949 172.23.3.135 -> 52.206.84.26 TCP 66 43074 > http [ACK] Seq=1 A
|||
|||
||| 9.942741717 172.23.3.135 -> 52.206.84.26 HTTP 113 GET / HTTP/1.1
||| 9.943058252 172.23.3.135 -> 52.206.84.26 TCP 66 43074 > http [FIN, ACK] Se
||| 9.943200815 52.206.84.26 -> 172.23.3.135 TCP 66 http > 43074 [ACK] Seq=1 A
||| 9.964304451 52.206.84.26 -> 172.23.3.135 HTTP 600 HTTP/1.1 302 Found (tex
||| 9.964317286 172.23.3.135 -> 52.206.84.26 TCP 66 43074 > http [ACK] Seq=49
||| 9.964319806 52.206.84.26 -> 172.23.3.135 TCP 66 http > 43074 [FIN, ACK] Se
||| 9.964322372 172.23.3.135 -> 52.206.84.26 TCP 66 43074 > http [ACK] Seq=49

```

# Resource: Network VIII

```
.....
.....
0.000 ( 0.000 ms): ... [continued]: select() = 1
686549238.395 ( 0.033 ms): read(buf: 0x1fc6040, count: 8192 ) ...
686549238.427 ( ) : probe:vfs_read:(ffffff811f9f40))
686549238.395 ( 0.061 ms): ... [continued]: read() = 47
686549238.483 ( 0.028 ms): write(fd: 3<socket:[354463]>, buf: 0x1fc6040, count: 47 ) ...
686549238.512 ( ) : probe:vfs_write:(ffffff811fa070))
686549238.543 ( ) : net:net_dev_queue:dev=eth0 skbaddr=0xffff8801c87550e8 len=113)
686549238.578 ( ) : net:net_dev_start_xmit:dev=eth0 queue_mapping=0 skbaddr=0xffff8801c87550e8 vlan_tagged=0 vlan_proto=0x0000
686549238.605 ( ) : probe:ixgbev_f_xmit_frame:(fffffffa01d9820))
686549238.633 ( ) : net:net_dev_xmit:dev=eth0 skbaddr=0xffff8801c87550e8 len=113 rc=0)
686549238.483 ( 0.179 ms): ... [continued]: write() = 47
686549238.692 ( 0.028 ms): select(n: 4, inp: 0x7ffe7ef65f60, outp: 0x7ffe7ef66060, exp: 0x7ffe7ef66060) = 2
686549238.747 ( 0.026 ms): read(buf: 0x1fc6040, count: 8192 ) ...
686549238.772 ( ) : probe:vfs_read:(ffffff811f9f40))
686549238.747 ( 0.052 ms): ... [continued]: read() = 0
686549238.828 ( 0.030 ms): shutdown(fd: 3<socket:[354463]>, how: 1 ) ...
686549238.859 ( ) : net:net_dev_queue:dev=eth0 skbaddr=0xffff8801c8756ae8 len=66)
686549238.894 ( ) : net:net_dev_start_xmit:dev=eth0 queue_mapping=0 skbaddr=0xffff8801c8756ae8 vlan_tagged=0 vlan_proto=0x0000
686549238.920 ( ) : probe:ixgbev_f_xmit_frame:(fffffffa01d9820))
686549238.946 ( ) : net:net_dev_xmit:dev=eth0 skbaddr=0xffff8801c8756ae8 len=66 rc=0)
686549238.828 ( 0.147 ms): ... [continued]: shutdown() = 0
686549239.003 ( 21.167 ms): select(n: 4, inp: 0x7ffe7ef65f60, outp: 0x7ffe7ef65f60, exp: 0x7ffe7ef66060, tvp: 0x7ffe7ef66170) = 1
686549260.213 ( 0.025 ms): read(fd: 3<socket:[354463]>, buf: 0x1fc6040, count: 8192 ) ...
686549260.239 ( ) : probe:vfs_read:(ffffff811f9f40))
686549260.213 ( 0.056 ms): ... [continued]: read() = 534
686549260.297 ( 0.027 ms): write(fd: 1</dev/pts/4>, buf: 0x1fc6040, count: 534 ) ...
686549260.324 ( ) : probe:vfs_write:(ffffff811fa070))
686549260.297 ( 0.065 ms): ... [continued]: write() = 534
686549260.389 ( 0.027 ms): shutdown(fd: 3<socket:[354463]>, how: 1 ) = -1 ENOTCONN Transport endpoint i
686549260.447 ( 0.029 ms): select(n: 4, inp: 0x7ffe7ef65f60, outp: 0x7ffe7ef65f60, exp: 0x7ffe7ef66060, tvp: 0x7ffe7ef66170) = 2
686549260.503 ( 0.027 ms): read(fd: 3<socket:[354463]>, buf: 0x1fc6040, count: 8192 ) ...
686549260.530 ( ) : probe:vfs_read:(ffffff811f9f40))
686549260.503 ( 0.053 ms): ... [continued]: read() = 0
686549260.584 ( 0.030 ms): shutdown(fd: 3<socket:[354463]>, how: 1 ) = -1 ENOTCONN Transport endpoint i
686549260.649 ( 0.029 ms): ioctl(fd: 1</dev/pts/4>, cmd: TCSETS, arg: 0x7ffe7ef65f70 ) = 0
686549260.704 ( 0.026 ms): ioctl(fd: 1</dev/pts/4>, cmd: TCGETS, arg: 0x7ffe7ef65f70 ) = 0
686549260.758 ( 0.026 ms): shutdown(fd: 3<socket:[354463]>, how: 2 ) = -1 ENOTCONN Transport endpoint i
686549260.881 ( 0.000 ms): exit_group( )
```

# Resource: Network IX

```
# Delete the probes
perf probe -m ixgbevfv --del='ixgbevfv_*'
perf probe --del='vfv_*'
```

## ● Utilization at system level

```
# Utilization
```

```
# nethogs
```

```
NetHogs version 0.8.5
```

PID	USER	PROGRAM	DEV	SENT	RECEIVED
1281	deploy	..ssenger RubyApp: /data/helpkit/current/public (prod	eth0	27.519	279.612 KB/sec
25051	deploy	..ssenger RubyApp: /data/helpkit/current/public (prod	eth0	33.842	279.514 KB/sec
3044	deploy	..ssenger RubyApp: /data/helpkit/current/public (prod	eth0	12.543	86.839 KB/sec
20292	deploy	..ssenger RubyApp: /data/helpkit/current/public (prod	eth0	36.001	86.133 KB/sec
21235	deploy	nginx: worker process	eth0	98.778	22.356 KB/sec
7154	root	/usr/bin/python	eth0	122.532	3.029 KB/sec
8073	deploy	..ssenger RubyApp: /data/helpkit/current/public (prod	eth0	0.801	2.462 KB/sec
9090	deploy	..ssenger RubyApp: /data/helpkit/current/public (prod	eth0	0.454	1.387 KB/sec
9410	root	tail	eth0	0.392	0.381 KB/sec
16816	root	/opt/SumoCollector/jre/bin/java	eth0	1.734	0.380 KB/sec
12620	suresh	sshd: suresh@pts/0	eth0	0.178	0.052 KB/sec
?	root	10.2.16.117:45404-54.231.141.84:443		0.000	0.000 KB/sec
?	root	10.2.16.117:53200-10.2.204.10:9101		0.000	0.000 KB/sec
?	root	unknown TCP		0.000	0.000 KB/sec

TOTAL 334.773 762.145 KB/sec



# Resource: Network X

#####

# iftop

		1.91Mb	3.81Mb	5.72Mb	7.63Mb	9.54Mb
rails-app-4.localdomain	=>	ip-10-2-86-132.eu-west-1.compute.i	460Kb	310Kb	310Kb	
	<=		8.64Mb	5.09Mb	5.09Mb	
rails-app-4.localdomain	=>	ip-10-2-86-150.eu-west-1.compute.i	429Kb	247Kb	247Kb	
	<=		5.77Mb	3.25Mb	3.25Mb	
rails-app-4.localdomain	=>	ec2-34-253-108-119.eu-west-1.compu	2.06Mb	1.48Mb	1.48Mb	
	<=		28.2Kb	25.2Kb	25.2Kb	
rails-app-4.localdomain	=>	ip-10-2-20-57.eu-west-1.compute.in	96.1Kb	77.0Kb	77.0Kb	
	<=		745Kb	807Kb	807Kb	
rails-app-4.localdomain	=>	ip-10-2-10-9.eu-west-1.compute.int	729Kb	519Kb	519Kb	
	<=		385Kb	255Kb	255Kb	
rails-app-4.localdomain	=>	ip-10-2-21-95.eu-west-1.compute.in	86.6Kb	79.1Kb	79.1Kb	
	<=		559Kb	682Kb	682Kb	
rails-app-4.localdomain	=>	collector-3.newrelic.com	0b	415Kb	415Kb	
	<=		0b	9.27Kb	9.27Kb	
rails-app-4.localdomain	=>	ip-10-2-10-48.eu-west-1.compute.in	156Kb	207Kb	207Kb	
	<=		121Kb	114Kb	114Kb	
rails-app-4.localdomain	=>	ip-10-2-87-122.eu-west-1.compute.i	7.23Kb	32.7Kb	32.7Kb	
	<=		43.0Kb	265Kb	265Kb	
rails-app-4.localdomain	=>	ip-10-2-87-246.eu-west-1.compute.i	24.6Kb	33.7Kb	33.7Kb	
	<=		191Kb	232Kb	232Kb	
rails-app-4.localdomain	=>	ip-10-2-20-217.eu-west-1.compute.i	50.8Kb	56.5Kb	56.5Kb	
	<=		81.7Kb	73.2Kb	73.2Kb	

# Resource: Network XI

TX:	cum:	1.77MB	peak:	4.23Mb	rates:	4.23Mb	3.55Mb	3.55Mb
RX:		5.41MB		16.6Mb		16.6Mb	10.8Mb	10.8Mb
TOTAL:		7.18MB		20.8Mb		20.8Mb	14.4Mb	14.4Mb

## ● View socket connections

```
# ss
Netid State Recv-Q Send-Q Local Address:Port Peer Address:Port
u_str ESTAB 0 0 * 655403 * 655402
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/preloader.1uya910 81482655
u_str ESTAB 0 0 * 81813975 * 81813976
u_str ESTAB 0 0 * 81803882 * 81804840
.....
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/preloader.1uya910 81804840
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/ruby.JYtJfWTrsC07A8epKaneYMZ58GHSsR6NAhkVUkCKaeqa
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/preloader.1uya910 81813976
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/ruby.JX11S6SGXlpBiU6QL12YgHjSHQwgmUHI8wdW0jXQB0nL
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/preloader.1uya910 81670404
u_str ESTAB 0 0 * 81876765 * 81874494
u_str ESTAB 0 0 /tmp/passenger.RgJUBdx/apps.s/ruby.htx8W1KaQ8RwQ6K1CzhIyElbVvBUS4oNR6BPrJmsbCrJ
tcp ESTAB 0 0 10.2.16.117:58804 50.31.164.148:http
tcp ESTAB 0 0 10.2.16.117:60228 10.2.86.42:mysql
tcp LAST-ACK 0 32 10.2.16.117:54172 52.94.5.156:https
tcp ESTAB 0 0 10.2.16.117:47040 10.2.20.217:6379
tcp ESTAB 0 0 10.2.16.117:45890 10.2.20.217:6379
.....
tcp ESTAB 0 0 10.2.16.117:59738 10.2.21.162:memcache
```

# Resource: Network XII

```
tcp    ESTAB      0      0      10.2.16.117:45704      10.2.20.217:6379
tcp    ESTAB      0      0      10.2.16.117:46888      10.2.20.64:memcache
.....
tcp    ESTAB      0      0      10.2.16.117:38134      50.31.164.149:http
tcp    ESTAB      0      0      10.2.16.117:45154      50.31.164.147:http
tcp    ESTAB      0      0      10.2.16.117:52644      10.2.204.15:http
tcp    ESTAB      0      0      10.2.16.117:41902      10.2.20.76:6379
tcp    ESTAB      0      0      10.2.16.117:45784      10.2.20.217:6379
tcp    CLOSE-WAIT 1      0      ::ffff:10.2.16.117:46754      ::ffff:169.254.169.254:http
tcp    CLOSE-WAIT 32     0      ::ffff:10.2.16.117:47502      ::ffff:176.34.227.36:https
tcp    ESTAB      0      0      ::ffff:127.0.0.1:31000      ::ffff:127.0.0.1:32000
tcp    ESTAB      0      0      ::ffff:10.2.16.117:55954      ::ffff:46.51.173.146:https
```

#####

You can get even more detailed information about a socket from kernel's internal socket struct. For example, we will try to get nginx listening (on port 81) socket's backlog length

NOTE: You may need to install kernel debug info if not already installed

```
# yum-config-manager --enable "amzn-main-debuginfo" --enable "amzn-updates-debuginfo"
```

```
# yum -y install kernel-debuginfo kernel-devel
```

Or get socket info for listening socket on port 81

```
# ss -ltn | grep :81
```

```
tcp    LISTEN      0      511      *:81          *:          ino:29842919 sk:55 <->
```

Get its sk buff address

```
# grep 29842919 /proc/net/tcp
```

```
8: 00000000:0051 00000000:0000 0A 00000000:00000000 00:00000000 00000000      0      0 29842919 1 ffff8800e
```

# Resource: Network XIII

```
Now get the details, for example, the backlog length
# gdb /usr/lib/debug/lib/modules/'uname -r'/vmlinux /proc/kcore
.....
.....
Reading symbols from /usr/lib/debug/lib/modules/4.4.51-40.69.amzn1.x86_64/vmlinux...done.
[New process 1]
Core was generated by 'root=LABEL=/ console=tty1 console=ttyS0 selinux=0 LANG=en_US.UTF-8 KEYTABLE=us'.
.....
(gdb) set print pretty on
(gdb) p *(struct sock *)0xffff8800e9d1da00
.....
.....
      sk_ack_backlog = 0,
      sk_max_ack_backlog = 511,
.....
.....
```

## • Network latency / reachability

```
# Latency
# ping -c 3 google.com
PING google.com (172.217.7.142) 56(84) bytes of data.
64 bytes from iad30s08-in-f14.1e100.net (172.217.7.142): icmp_seq=1 ttl=48 time=1.57 ms
64 bytes from iad30s08-in-f142.1e100.net (172.217.7.142): icmp_seq=2 ttl=48 time=1.12 ms
64 bytes from iad30s08-in-f14.1e100.net (172.217.7.142): icmp_seq=3 ttl=48 time=1.13 ms

--- google.com ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
```

- Errors

# Look for errors, dropped, overruns etc

# ethtool -S eth0

NIC statistics:

```
rx_packets: 604359552
tx_packets: 649217118
rx_bytes: 616073831008
tx_bytes: 278008238445
tx_busy: 0
tx_restart_queue: 0
tx_timeout_count: 0
multicast: 0
rx_csum_offload_errors: 0
rx_bp_poll_yield: 0
rx_bp_cleaned: 0
rx_bp_misses: 0
tx_bp_napi_yield: 0
tx_bp_cleaned: 0
tx_bp_misses: 0
```

- Looking at live network traffic

# Resource: Network XV

Just display the TCP connection establishment alone

```
# tshark -f '(tcp[tcpflags] & (tcp-syn) != 0)'
```

Running as user "root" and group "root". This could be dangerous.

Capturing on eth0

```
0.000000000 172.16.10.27 -> 172.16.17.173 TCP 74 40478 > 81 [SYN] Seq=0 Win=26883 Len=0 MSS=8961 SACK_PERM=1 TS
0.000015819 172.16.17.173 -> 172.16.10.27 TCP 74 81 > 40478 [SYN, ACK] Seq=0 Ack=1 Win=26847 Len=0 MSS=8961 SAC
0.063601284 172.16.10.153 -> 172.16.17.173 TCP 74 56494 > 81 [SYN] Seq=0 Win=26883 Len=0 MSS=8961 SACK_PERM=1 T
0.063618642 172.16.17.173 -> 172.16.10.153 TCP 74 81 > 56494 [SYN, ACK] Seq=0 Ack=1 Win=26847 Len=0 MSS=8961 SA
0.074333351 172.16.10.13 -> 172.16.17.173 TCP 74 37818 > 81 [SYN] Seq=0 Win=26883 Len=0 MSS=8961 SACK_PERM=1 TS
.....
.....
```

# Resource: Various software resources I

- Global resource limits
- Process/thread specific resource limits
- cgroup
- Lock contention

# Certain resource limits may be set at system level as well as process level

# Maximum system level file descriptors

# sysctl fs.file-max fs.file-nr

fs.file-max = 762054

fs.file-nr = 1888            0            762054

# Maximum system level processes/threads

# sysctl kernel.threads-max kernel.pid\_max

kernel.threads-max = 59690

kernel.pid\_max = 32768

# Process level limits

# cat /proc/\$(pidof ruby | awk '{print \$1}')

Limit	Soft Limit	Hard Limit	Units
Max cpu time	unlimited	unlimited	seconds
Max file size	unlimited	unlimited	bytes
Max data size	unlimited	unlimited	bytes
Max stack size	8388608	unlimited	bytes

# Resource: Various software resources II

Max core file size	0	unlimited	bytes
Max resident set	unlimited	unlimited	bytes
Max processes	29845	29845	processes
Max open files	1024	4096	files
Max locked memory	65536	65536	bytes
Max address space	unlimited	unlimited	bytes
Max file locks	unlimited	unlimited	locks
Max pending signals	29845	29845	signals
Max msgqueue size	819200	819200	bytes
Max nice priority	0	0	
Max realtime priority	0	0	
Max realtime timeout	unlimited	unlimited	us

```
# ls -l /proc/$(pidof ruby | awk '{print $1}')/fd | wc -l
23
```

```
# Number of threads in a given process
# ls -l /proc/$(pidof ruby | awk '{print $1}')/task/ | wc -l
6
```

```
# Where/on what a given process is waiting on
# cat /proc/$(pidof ruby | awk '{print $1}')/stack
[<ffffff811e8219>] poll_schedule_timeout+0x49/0x70
[<ffffff811e8bac>] do_select+0x58c/0x750
[<ffffff811e8f3c>] core_sys_select+0x1cc/0x2d0
[<ffffff811e90eb>] Sys_select+0xab/0xf0
[<ffffff814f002e>] entry_SYSCALL_64_fastpath+0x12/0x71
[<ffffffffffffffff>] 0xffffffffffffffff
```



# Xen specific: Exploring xenstore from DomU I

```
# Install dependencies
$ yum groupinstall -y 'Development Tools'
$ yum install -y dev86 iasl ncurses-devel glib2-devel pixman-devel \
libaio-devel glibc-devel.i686 cmake xz-devel libuuid-devel \
zlib-devel
$ pushd ~
$ wget http://github.com/lloyd/yajl/tarball/2.1.0
$ tar xvf 2.1.0
$ cd lloyd-yajl-66cb08c/
$ ./configure
$ make
$ sudo make install
$ popd

# Install xen tools
$ wget https://downloads.xenproject.org/release/xen/4.9.0/xen-4.9.0.tar.gz
$ tar xvf xen-4.9.0.tar.gz
$ cd xen-4.9.0/tools/
$ ./configure
$ make -C include
$ make -C ./libs
$ make -C ./libxc
$ make -C ./xenstore
$ sudo make install -C ./xenstore bindir=/usr/local/bin libdir=/usr/local/lib
$ export PATH=$PATH:/usr/local/bin
$ export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/lib
$ sudo ldconfig

$ sudo mount -t xenfs none /proc/xen
```

# Xen specific: Exploring xenstore from DomU II

```
#####
```

```
# List the store exposed to this DomU
$ sudo su
$ xenstore-ls /local/domain/'xenstore-read domid'
vm = "/vm/ec2a1431-14c4-fed6-dabd-eed158b16aa4"
device = ""
vbd = ""
51712 = ""
    backend-id = "0"
    virtual-device = "51712"
    device-type = "disk"
    state = "4"
    backend = "/local/domain/0/backend/vbd/204/51712"
    ring-ref = "8"
    event-channel = "35"
    protocol = "x86_64-abi"
    feature-persistent = "1"
51824 = ""
    backend-id = "0"
    virtual-device = "51824"
    device-type = "disk"
    state = "4"
    backend = "/local/domain/0/backend/vbd/204/51824"
    ring-ref = "1090"
    event-channel = "40"
    protocol = "x86_64-abi"
    feature-persistent = "1"
pci = ""
```

# Xen specific: Exploring xenstore from DomU III

```
0 = ""
  state = "1"
  backend-id = "0"
  backend = "/local/domain/0/backend/pci/204/0"
console = ""
0 = ""
  state = "1"
  backend-id = "0"
  backend = "/local/domain/0/backend/console/204/0"
control = ""
platform-feature-multiprocessor-suspend = "1"
platform-feature-xs_reset_watches = "1"
error = ""
memory = ""
  target = "7864320"
guest = ""
hvmprv = ""
data = ""
image = ""
  device-model-fifo = "/var/run/xend/dm-204-1502179628.fifo"
  device-model-pid = "11313"
  suspend-cancel = "1"
console = ""
vnc-port = "5905"
vnc-listen = "127.0.0.1"
vnc-pass = "xyz"
port = "7"
limit = "1048576"
type = "ioemu"
serial = ""
```

# Xen specific: Exploring xenstore from DomU IV

```
0 = ""
  tty = "/dev/pts/5"
description = ""
cpu = ""
2 = ""
  availability = "online"
0 = ""
  availability = "online"
3 = ""
  availability = "online"
1 = ""
  availability = "online"
domid = "204"
store = ""
  ring-ref = "1044476"
  port = "6"
name = "dom_27677293465"
device-misc = ""
  console = ""
  nextDeviceID = "1"
```

```
# List one specific VBD device from above
$ xenstore-ls /local/domain/0/backend/vbd/826/51728
domain = "dom_24698651860"
frontend = "/local/domain/826/device/vbd/51728"
uuid = "d62daa8e-d864-e843-f455-3640ffc3cbbf"
bootable = "0"
dev = "xvdb"
state = "4"
params = "/dev/nvme2n1"
```

# Xen specific: Exploring xenstore from DomU V

```
mode = "w"  
removable = "1"  
online = "1"  
frontend-id = "826"  
type = "phy"  
physical-device = "fb:40"  
hotplug-status = "connected"  
feature-flush-cache = "0"  
feature-discard = "0"  
feature-barrier = "0"  
feature-persistent = "1"  
feature-max-indirect-segments = "256"  
sectors = "209715200"  
info = "0"  
sector-size = "512"  
physical-sector-size = "512"
```

- 1 System
- 2 Application**
- 3 Profilers
- 4 Debuggers
- 5 Books

# Application performance analysis


- Resource based USE method is what we saw so far for identifying system level performance issues.
- For identifying application performance issues, we could identify what the application is doing over a period then analyze: TSA (Thread State Analysis) method <sup>7</sup> is something we could use.
- For the application/process we want to analyze, identify the threads it has
- For each thread
  - Measure time spent in each state
    - State can be R (running), S (sleeping), D (uninterruptible sleep), T (stopped), t (stopped by debugger), Z (zombie)
  - Investigate states from most frequent to least
- But before we introduce/use various tools to do that, we have to know various things about process, so we will take a detour and do a deep dive into process

---


<sup>7</sup><http://www.brendangregg.com/tsamethod.html>

- Introduction: <http://duartes.org/gustavo/blog/post/anatomy-of-a-program-in-memory/>
- User Space vs Kernel Space split (32bit OS)


Linux User/Kernel  
Memory Split



Windows, default  
memory split




Windows booted  
with /3GB switch





# Process II

- Kernel Space is same across processes, only User Space content changes across processes



- Process is always started from an executable binary file in certain format (ELF is the most common in Linux). That format defines/standardizes various things that would be read/used by OS during process creation
  - A typical example, for Golang: Your source code is compiled into native machine code and an ELF file is created with all the info, including the generated machine code, that OS can use during process creation.
  - Note that scripted applications, like Ruby, first a process (VM, generally written in C/C++) is started, which has logic to parse/run the ruby scripts further.

- The executable binary may be
  - Statically linked: All dependent code/data is included in the binary and is self contained
  - Or dynamically linked: Certain code/data it depends on comes from another binary (typically a shared library .so) and will only be resolved during process startup. The tool that does "runtime resolving" is called a "loader".
- Let's take a look at how the executable looks on disk.  
We will use ruby executable as an example

- What kind of file it is?

```
$ file /usr/local/bin/ruby
/usr/local/bin/ruby: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked,
interpreter /lib64/ld-linux-x86-64.so.2, for GNU/Linux 2.6.35,
BuildID[sha1]=27875858789fb14bfbf4ac2c603ec700acff91da, not stripped
```

It is 64 bit ELF binary, dynamically linked, using loader at /lib64/ld-linux-x86-64.so.2 and contains debug symbols

- What it depends on (shared libraries)?

```
$ ldd /usr/local/bin/ruby
linux-vdso.so.1 => (0x00007ffe9c99c000)
libruby.so.2.2 => /usr/local/lib/libruby.so.2.2 (0x00007fbb49095000)
libpthread.so.0 => /lib64/libpthread.so.0 (0x00007fbb48e79000)
libdl.so.2 => /lib64/libdl.so.2 (0x00007fbb48c74000)
libcrypt.so.1 => /lib64/libcrypt.so.1 (0x00007fbb48a3d000)
libm.so.6 => /lib64/libm.so.6 (0x00007fbb4873b000)
libc.so.6 => /lib64/libc.so.6 (0x00007fbb48378000)
/lib64/ld-linux-x86-64.so.2 (0x00005634741af000)
libfreebl3.so => /lib64/libfreebl3.so (0x00007fbb48176000)
```

- What it contains?

```
$ size /usr/local/bin/ruby
   text    data     bss       dec      hex filename
   2298     668        4      2970     b9a /usr/local/bin/ruby
```

```
# Or all sections
$ size -Ax /usr/local/bin/ruby
/usr/local/bin/ruby :
section              size      addr
.interp              0x1c     0x400200
.note.ABI-tag        0x20     0x40021c
.note.gnu.build-id   0x24     0x40023c
.gnu.hash            0x50     0x400260
.dynsym              0x228    0x4002b0
.dynstr              0x164    0x4004d8
.gnu.version         0x2e     0x40063c
.gnu.version_r       0x20     0x400670
.rela.dyn            0x18     0x400690
```

# Process V

.rela.plt	0xc0	0x4006a8
.init	0x1a	0x400768
.plt	0x90	0x400790
.text	0x1b4	0x400820
.fini	0x9	0x4009d4
.rodata	0x11	0x4009e0
.eh_frame_hdr	0x34	0x4009f4
.eh_frame	0xec	0x400a28
.init_array	0x8	0x600b18
.fini_array	0x8	0x600b20
.jcr	0x8	0x600b28
.dynamic	0x220	0x600b30
.got	0x8	0x600d50
.got.plt	0x58	0x600d58
.data	0x4	0x600db0
.bss	0x4	0x600db4
.comment	0x2c	0x0
.debug_aranges	0x30	0x0
.debug_info	0x492	0x0
.debug_abbrev	0x18d	0x0
.debug_line	0x538	0x0
.debug_str	0x194f1	0x0
.debug_loc	0xaa	0x0
.debug_ranges	0x20	0x0
.debug_macro	0x56ac	0x0
Total	0x203b4	

- How does ELF format look like?

# Process VI

```
# Show the header
$ readelf -h /usr/local/bin/ruby
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                               2's complement, little endian
  Version:                             1 (current)
  OS/ABI:                              UNIX - System V
  ABI Version:                         0
  Type:                                EXEC (Executable file)
  Machine:                             Advanced Micro Devices X86-64
  Version:                             0x1
  Entry point address:                 0x400870
  Start of program headers:            64 (bytes into file)
  Start of section headers:            132928 (bytes into file)
  Flags:                               0x0
  Size of this header:                 64 (bytes)
  Size of program headers:             56 (bytes)
  Number of program headers:           8
  Size of section headers:             64 (bytes)
  Number of section headers:           38
  Section header string table index:   35
```

```
# Show various sections
$ readelf -S /usr/local/bin/ruby
There are 38 section headers, starting at offset 0x20740:
```

Section Headers:

[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info	Align

# Process VII

```
[ 0]          NULL          0000000000000000 00000000
    0000000000000000 0000000000000000      0      0      0
[ 1] .interp      PROGBITS      0000000000400200 00000200
    000000000000001c 0000000000000000      A      0      0      1
.....
.....
[13] .text          PROGBITS      0000000000400820 00000820
    00000000000001b4 0000000000000000      AX      0      0      16
[14] .fini          PROGBITS      00000000004009d4 000009d4
    0000000000000009 0000000000000000      AX      0      0      4
[15] .rodata        PROGBITS      00000000004009e0 000009e0
    0000000000000011 0000000000000000      A      0      0      8
.....
.....
[28] .debug_info     PROGBITS      0000000000000000 00000e10
    0000000000000492 0000000000000000      0      0      1
[29] .debug_abbrev   PROGBITS      0000000000000000 000012a2
    000000000000018d 0000000000000000      0      0      1
.....
.....
```

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), l (large)  
I (info), L (link order), G (group), T (TLS), E (exclude), x (unknown)  
0 (extra OS processing required) o (OS specific), p (processor specific)

```
# Show the segment header (i.e., on memory layout)
$ readelf -l 'which ruby'
```

Elf file type is EXEC (Executable file)

# Process VIII

Entry point 0x400870

There are 8 program headers, starting at offset 64

Program Headers:

Type	Offset	VirtAddr	PhysAddr
	FileSiz	MemSiz	Flags Align
PHDR	0x0000000000000040	0x0000000000400040	0x0000000000400040
	0x00000000000001c0	0x00000000000001c0	R E 8
INTERP	0x0000000000000200	0x0000000000400200	0x0000000000400200
	0x00000000000001c	0x00000000000001c	R 1
[Requesting program interpreter: /lib64/ld-linux-x86-64.so.2]			
LOAD	0x0000000000000000	0x0000000000400000	0x0000000000400000
	0x0000000000000b14	0x0000000000000b14	R E 200000
LOAD	0x0000000000000b18	0x0000000000600b18	0x0000000000600b18
	0x000000000000029c	0x00000000000002a0	RW 200000
DYNAMIC	0x0000000000000b30	0x0000000000600b30	0x0000000000600b30
	0x0000000000000220	0x0000000000000220	RW 8
NOTE	0x000000000000021c	0x000000000040021c	0x000000000040021c
	0x0000000000000044	0x0000000000000044	R 4
GNU_EH_FRAME	0x00000000000009f4	0x00000000004009f4	0x00000000004009f4
	0x0000000000000034	0x0000000000000034	R 4
GNU_STACK	0x0000000000000000	0x0000000000000000	0x0000000000000000
	0x0000000000000000	0x0000000000000000	RW 10

Section to Segment mapping:

Segment Sections...

00

01 .interp

02 .interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version\_r

03 .init\_array .fini\_array .jcr .dynamic .got .got.plt .data .bss

# Process IX

```
04      .dynamic
05      .note.ABI-tag .note.gnu.build-id
06      .eh_frame_hdr
07
```

```
# Or use objdump to see them together
$ objdump -h /usr/local/bin/ruby
```

```
/usr/local/bin/ruby:      file format elf64-x86-64
```

Sections:

Idx	Name	Size	VMA	LMA	File off	Algn
0	.interp	0000001c	0000000000400200	0000000000400200	00000200	2**0
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
.....						
11	.plt	00000090	0000000000400790	0000000000400790	00000790	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
12	.text	000001b4	0000000000400820	0000000000400820	00000820	2**4
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
13	.fini	00000009	00000000004009d4	00000000004009d4	000009d4	2**2
	CONTENTS, ALLOC, LOAD, READONLY, CODE					
14	.rodata	00000011	00000000004009e0	00000000004009e0	000009e0	2**3
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
.....						
23	.data	00000004	0000000000600db0	0000000000600db0	00000db0	2**2
	CONTENTS, ALLOC, LOAD, DATA					
24	.bss	00000004	0000000000600db4	0000000000600db4	00000db4	2**2
	ALLOC					



```
.....  
.....  
25 .comment      0000002c  0000000000000000  0000000000000000  00000db4  2**0  
                  CONTENTS, READONLY  
26 .debug_aranges 00000030  0000000000000000  0000000000000000  00000de0  2**0  
                  CONTENTS, READONLY, DEBUGGING  
.....  
.....  
33 .debug_macro   000056ac  0000000000000000  0000000000000000  0001af22  2**0  
                  CONTENTS, READONLY, DEBUGGING
```

- Where is my function in it?

```
# Use nm to look for symbols, for example, the entry point 0x400870 we found above,  
# we can check what function is contained there.
```

```
$ nm /usr/local/bin/ruby  
0000000000600db4 B __bss_start  
.....  
.....  
0000000000400820 t main  
00000000004008d0 t register_tm_clones  
                  U ruby_init  
                  U ruby_init_stack  
                  U ruby_options  
                  U ruby_run_node  
                  U ruby_sysinit  
                  U setlocale@@GLIBC_2.2.5  
0000000000400870 T _start  
.....  
.....
```

- How a process is started

You can find a detailed explanation here <sup>8</sup>

- Some process wants to start a new program (say, shell, wants to start /usr/local/bin/ruby program), so it calls into kernel (using fork/exec)
- Kernel checks what type of file it is (binfmt kernel feature)
  - It could be ELF, java binary, .net binary, shell script with shebang "#!/bin/sh" etc
  - Linux has extensible support via binfmt <sup>9</sup>
- If it is a supported file, in this case ELF executable, it will load it and pass control to <sup>10</sup>
  - "Load" == finding various ELF sections we saw above, memory mapping them
  - "Pass" == passing execution control to the entry point specified in the ELF binary
- Additionally what executable gets loaded/run may vary based whether the executable is statically linked or dynamically linked
- If it is dynamically linked

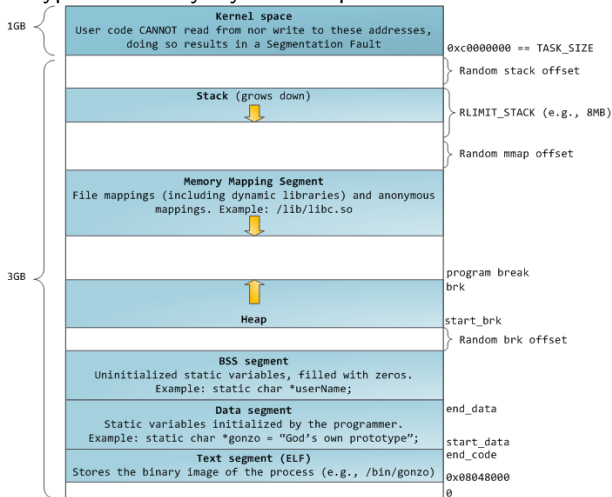
- Instead of loading the executable, it will load the loader (/lib64/ld-linux-x86-64.so.2) and pass control to it, along with info about the file to be executed (via AUX info). Note that loader is just another normal executable as far as kernel is concerned.

```
$ LD_SHOW_AUXV=1 /usr/local/bin/ruby
AT_SYSINFO_EHDR: 0x7ffdf92cf000
AT_HWCAP:        178bfbff
AT_PAGESZ:       4096
AT_CLKTCK:       100
AT_PHDR:         0x400040
AT_PENT:         56
AT_PHNUM:        8
AT_BASE:         0x7f32eb690000
AT_FLAGS:        0x0
AT_ENTRY:        0x400870
AT_UID:          2147
AT_EUID:         2147
AT_GID:          501
AT_EGID:         501
AT_SECURE:       0
AT_RANDOM:       0x7ffdf928b049
AT_EXECFN:       /usr/local/bin/ruby
AT_PLATFORM:     x86_64
```

- Loader in turn loads the /usr/local/bin/ruby executable and passes control to the entry point specified

- If it is statically linked, kernel loads the executable and passes control to the entry point specified
- Note that in either case (statically or dynamically linked), as far as kernel is concerned, it is just going to load an executable and pass control to it. In the case of dynamically linked executable, it just happens to be the "loader" executable.
- Additionally, other executables, like Java etc can be started via similar method (i.e., kernel loads the JVM and passes control to it etc)
- I have a detailed info here that covers `execve` syscall + loader initialization + symbol resolving + binary execution
- Let's take a look at how it looks at runtime

- A typical memory layout of a process



- How do I see a process's memory layout?

# Process XV

```
cat /proc/19206/maps
00400000-00401000 r-xp 00000000 ca:01 29854 /usr/local/bin/ruby
00600000-00601000 rw-p 00000000 ca:01 29854 /usr/local/bin/ruby
01cca000-1f6ca000 rw-p 00000000 00:00 0 [heap]
1f6ca000-4123c000 rw-p 00000000 00:00 0 [heap]
7f21ef922000-7f21ef923000 ---p 00000000 00:00 0
7f21ef923000-7f21efa23000 rw-p 00000000 00:00 0
7f21efa23000-7f21efa24000 ---p 00000000 00:00 0
7f21efa24000-7f21efb24000 rw-p 00000000 00:00 0
7f21efb24000-7f21efb44000 r-xp 00000000 ca:01 3880 /usr/lib64/libnssdbm3.so
.....
.....
7f21f10f7000-7f21f10fb000 r-xp 00000000 ca:70 524440 /data/helpkit/shared/bun
7f21fc0da000-7f21fc2d9000 ---p 00258000 ca:01 285180 /usr/lib64/mysql/libmysq
.....
.....
7f2209462000-7f2209726000 r-xp 00000000 ca:01 29858 /usr/local/lib/libruby.s
7f2209726000-7f2209925000 ---p 002c4000 ca:01 29858 /usr/local/lib/libruby.s
.....
.....
7f2209baf000-7f2209bcf000 r-xp 00000000 ca:01 268462 /lib64/ld-2.17.so
.....
.....
7ffffb3d23000-7ffffb4522000 rw-p 00000000 00:00 0 [stack]
7ffffb4537000-7ffffb4539000 r--p 00000000 00:00 0 [vvar]
7ffffb4539000-7ffffb453b000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 r-xp 00000000 00:00 0 [vsyscall]
```

- Stack

# Process XVI

- Heap
  - Where all the dynamic variables/objects from the process are stored
  - How it is managed: Manually or automatically

- Virtual vs Resident

- Resident: RSS vs PSS vs USS

```
# Check memory usage, system-wide
```

```
$ sudo smem -k
```

PID	User	Command	Swap	USS	PSS	RSS
2972	root	/sbin/mingetty /dev/tty4	0	88.0K	110.0K	1.4M
.....						
22190	deploy	Passenger AppPreloader: /da	0	141.4M	196.6M	509.4M
22800	deploy	Passenger RubyApp: /data/he	0	160.8M	213.4M	520.7M
22788	deploy	Passenger RubyApp: /data/he	0	161.1M	213.5M	520.6M
22772	deploy	Passenger RubyApp: /data/he	0	161.4M	213.8M	520.9M
22744	deploy	Passenger RubyApp: /data/he	0	177.1M	226.9M	526.0M
22754	deploy	Passenger RubyApp: /data/he	0	177.1M	226.9M	526.3M
18599	root	/opt/SumoCollector/jre/bin/	0	233.1M	233.2M	235.3M
22731	deploy	Passenger RubyApp: /data/he	0	188.4M	237.3M	526.6M

- 'Thread': Heap is shared, stack is unique to each 'thread'
- Process state (R, S, D etc)
- Tools: top, htop, pmap, smem etc

<sup>8</sup>[https://github.com/0xAX/linux-insides/blob/master/Misc/program\\_startup.md](https://github.com/0xAX/linux-insides/blob/master/Misc/program_startup.md)

<sup>9</sup>[https://en.wikipedia.org/wiki/Binfmt\\_misc](https://en.wikipedia.org/wiki/Binfmt_misc)

<sup>10</sup>[https://github.com/torvalds/linux/blob/v4.13/fs/binfmt\\_elf.c#L679](https://github.com/torvalds/linux/blob/v4.13/fs/binfmt_elf.c#L679)

# Process runtime I

- Some processes may have a runtime and some may not have, based on what kind language they were built with
  - Minimal or no runtime: C/C++, Rust etc
  - With runtime: Java/JVM, C#/CLR, Go, Ruby, Python etc
- How they execute code
  - Compiled to native: C/C++, Go, Rust, C# + ngen (AOT), Java + AOT etc
  - Interpreted: MRI Ruby, Python, Perl, Erlang/BeamVM, Node.js, Java etc
    - Some of them may have intermediate form, but they can still be interpreted: Example, Java => bytecode => Interpreted
  - JITed: C#, Java, Node.js
    - Some may alternate between interpreted mode and JIT (example: Java, node.js)
    - Some always start in JITed mode: CLR/C#
- By how they manage memory <sup>11</sup>
  - Manual memory management: C/C++ etc



# Process runtime II

- Automatic memory management
  - Garbage collection: Java, Go, Ruby, Python etc
  - Reference counting: Objective-C, Python, Rust, C++
  - Resource Acquisition Is Initialization (RAII): Rust, C++
- By how they manage concurrency/parallelism (== threading)
  - Single threading
  - Multi threading
    - 1:1 threading: C, C++, Java, Ruby, Python etc
    - N:1 thread: Ruby fibers <sup>12</sup> , <sup>13</sup>
    - M:N threading: Golang, BeamVM (Erlang, Elixir)
  - Multi-process model
  - Evented vs Threaded
    - How blocking operations are handled
    - Nodejs, Go example
- Examples:
  - Single threaded: MRI Ruby
  - Evented: nginx, haproxy, nodejs, Ruby + EventMachine

# Process runtime III

- Evented + multi-threaded: go, nodejs: blocking operations are sent to thread pool
- Multi-process + evented: nginx, haproxy

---

<sup>11</sup><https://www.cs.virginia.edu/~cs415/reading/bacon-garbage.pdf>

<sup>12</sup>[http://schmurfy.github.io/2011/09/25/on\\_fibers\\_and\\_threads.html](http://schmurfy.github.io/2011/09/25/on_fibers_and_threads.html)

<sup>13</sup><http://oldmoe.blogspot.in/2008/08/ruby-fibers-vs-ruby-threads.html> 🔍 🔍 🔍

# System Calls I

- Introduction to syscalls <sup>14</sup> , <sup>15</sup>
- What is syscall: User Mode code requesting a service from Kernel Mode. Example: Writing to a file, sending a data out over network etc

**Table:** List of typical syscalls

syscall	what it does
open	Open a file, returns file descriptor
socket	Open a socket, returns file descriptor
read	Read data from a file descriptor (file, socket)
write	Write data to a file descriptor (file, socket)
close	Close a file descriptor
fork	Create a new process (out of current process)
exec/execve	Replace current process with new program
connect	Connect to remote host
accept	Accept a new connection on a socket
stat	Get file status
ioctl	Perform control functions on file descriptor
mmap	Map a file to the process address space
brk	Extend the heap pointer

# System Calls II

- Blocking vs Non-Blocking
  - Not all syscalls have non-blocking option
- Since syscall is the primary way processes interact with system, finding out what kind of syscalls a process is making could give us very good insight into what it is doing: This can be used for debugging, profiling or just for general understanding of a given process.
- How to find what kind of syscalls a process is making?: 'strace' or 'perf trace'
- How strace is implemented
  - Using ptrace interface <sup>16</sup>
  - This will cause two context switches for each syscall traced, can slow down the program significantly if it uses too many syscalls.
  - 'perf trace' is better in terms of performance, use that over strace if possible
- Things to remember

# System Calls III

- Performance impact of tracing in production: don't use it unless really required
- Underlying language 'runtime' semantics (evented, threaded: 1:1 threading or M:N threading model etc)
- Examples:

```
# Summarize system calls made
# strace -c ruby -e 'puts "hello world"'
hello world
```

% time	seconds	usecs/call	calls	errors	syscall
19.30	0.000183	1	191	97	open
15.61	0.000148	0	408		lstat
14.35	0.000136	1	115		read
6.96	0.000066	2	42		brk
5.91	0.000056	2	32		mmap
5.80	0.000055	1	97		fstat
5.38	0.000051	1	42		fcntl
4.11	0.000039	0	96		close
2.85	0.000027	1	44	41	ioctl
2.53	0.000024	1	35		geteuid
2.11	0.000020	1	23		mprotect
2.00	0.000019	1	34		getuid
2.00	0.000019	1	35		getegid
1.90	0.000018	1	18		rt_sigaction
1.69	0.000016	16	1		clone

# System Calls IV

1.69	0.000016	0	34	getgid
1.37	0.000013	7	2	pipe2
0.84	0.000008	3	3	getpid
0.63	0.000006	2	3	prlimit64
0.63	0.000006	3	2	getrandom
0.42	0.000004	0	31	4 stat
0.42	0.000004	2	2	futex
0.42	0.000004	4	1	sched_getaffinity
0.42	0.000004	4	1	clock_gettime
0.32	0.000003	1	3	rt_sigprocmask
0.21	0.000002	2	1	sigaltstack
0.11	0.000001	0	8	lseek
0.00	0.000000	0	2	write
0.00	0.000000	0	3	munmap
0.00	0.000000	0	1	1 access
0.00	0.000000	0	1	execve
0.00	0.000000	0	6	getdents
0.00	0.000000	0	1	arch_prctl
0.00	0.000000	0	1	set_tid_address
0.00	0.000000	0	1	set_robust_list
-----				
100.00	0.000948		1320	143 total

#####

```
# Or use 'perf trace'
# perf trace -s ruby -e 'puts "hello world"'
hello world
```

Summary of events:

# System Calls V

ruby (18436), 10 events, 0.4%

syscall	calls	total (msec)	min (msec)	avg (msec)	max (msec)	stddev (%)
read	2	0.004	0.002	0.002	0.003	21.82%
close	2	0.004	0.002	0.002	0.002	11.40%
poll	1	0.000	0.000	0.000	0.000	0.00%

ruby (18435), 2638 events, 98.7%

syscall	calls	total (msec)	min (msec)	avg (msec)	max (msec)	stddev (%)
lstat	408	1.215	0.002	0.003	0.009	1.69%
open	191	1.001	0.002	0.005	0.017	4.52%
read	115	0.620	0.002	0.005	0.040	7.77%
clone	1	0.315	0.315	0.315	0.315	0.00%
brk	42	0.270	0.001	0.006	0.014	7.19%
stat	31	0.177	0.002	0.006	0.010	8.57%
close	96	0.173	0.001	0.002	0.004	2.60%
fstat	97	0.170	0.001	0.002	0.003	2.50%
mmap	32	0.143	0.002	0.004	0.009	5.93%
mprotect	23	0.111	0.003	0.005	0.008	6.21%
ioctl	44	0.081	0.001	0.002	0.007	6.61%
fcntl	42	0.068	0.001	0.002	0.003	3.13%
futex	2	0.056	0.003	0.028	0.053	89.36%
getuid	34	0.053	0.001	0.002	0.004	5.53%

# System Calls VI

geteuid	35	0.048	0.001	0.001	0.004	5.22%
getegid	35	0.046	0.001	0.001	0.002	2.88%
getgid	34	0.045	0.001	0.001	0.002	3.10%
getdents	6	0.041	0.001	0.007	0.013	31.37%
munmap	3	0.032	0.007	0.011	0.013	16.94%
rt_sigaction	18	0.025	0.001	0.001	0.002	4.22%
lseek	8	0.014	0.001	0.002	0.002	6.17%
write	2	0.011	0.002	0.006	0.009	57.99%
pipe2	2	0.011	0.003	0.005	0.008	51.70%
getrandom	2	0.007	0.003	0.003	0.004	22.68%
rt_sigprocmask	3	0.005	0.001	0.002	0.003	26.26%
access	1	0.005	0.005	0.005	0.005	0.00%
prlimit64	3	0.004	0.001	0.001	0.001	3.48%
getpid	3	0.004	0.001	0.001	0.001	0.35%
clock_gettime	1	0.003	0.003	0.003	0.003	0.00%
sched_getaffinity	1	0.002	0.002	0.002	0.002	0.00%
arch_prctl	1	0.002	0.002	0.002	0.002	0.00%
sigaltstack	1	0.001	0.001	0.001	0.001	0.00%
set_robust_list	1	0.001	0.001	0.001	0.001	0.00%
set_tid_address	1	0.001	0.001	0.001	0.001	0.00%

#####

# Trace all syscalls

strace -ttT -ff ruby -e 'puts "hello world"'

```
11:15:42.151594 execve("/home/suresh/.rvm/rubies/ruby-2.4.1/bin/ruby", ["ruby", "-e", "puts \"hello world\"", [/* 119 vars */]) = 0 <0.000155>
11:15:42.151909 brk(NULL) = 0x2214000 <0.000007>
11:15:42.151955 access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory) <0.000009>
```



# System Calls VII

```
11:15:42.152005 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/x86_64/libruby.so.2.4", O_RDONLY|O_CLOEXEC)
= -1 ENOENT (No such file or directory) <0.000015>
11:15:42.152047 stat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/x86_64", 0x7fff6893f510) = -1 ENOENT
(No such file or directory) <0.000009>
11:15:42.152080 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/libruby.so.2.4", O_RDONLY|O_CLOEXEC)
= -1 ENOENT (No such file or directory) <0.000009>
11:15:42.152113 stat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls", 0x7fff6893f510) = -1 ENOENT (No such
file or directory) <0.000010>
11:15:42.152148 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/x86_64/libruby.so.2.4", O_RDONLY|O_CLOEXEC)
= -1 ENOENT (No such file or directory) <0.000010>
11:15:42.152181 stat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/x86_64", 0x7fff6893f510) = -1 ENOENT
(No such file or directory) <0.000015>
11:15:42.152234 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libruby.so.2.4", O_RDONLY|O_CLOEXEC) =
3 <0.000016>
11:15:42.152277 read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\360\335\2\0\0\0\0"..., 832)
= 832 <0.000009>
11:15:42.152318 fstat(3, {st_mode=S_IFREG|0755, st_size=4794408, ...}) = 0 <0.000008>
11:15:42.152349 mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f5124a2c000
<0.000009>
11:15:42.152380 mmap(NULL, 5339072, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f51242f3000
<0.000007>
11:15:42.152404 mprotect(0x7f51245f0000, 2097152, PROT_NONE) = 0 <0.000010>
11:15:42.152451 mmap(0x7f51247f0000, 32768, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE,
3, 0x2fd000) = 0x7f51247f0000 <0.000016>
11:15:42.152498 mmap(0x7f51247f8000, 75712, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS,
-1, 0) = 0x7f51247f8000 <0.000011>
11:15:42.152544 close(3) = 0 <0.000006>
11:15:42.152587 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libpthread.so.0", O_RDONLY|O_CLOEXEC) =
-1 ENOENT (No such file or directory) <0.000012>
.....
```

# System Calls VIII

```
11:15:42.166524 clone(strace: Process 5106 attached
child_stack=0x7f5124a2aff0,
    flags=CLONE_VM|CLONE_FS|CLONE_FILES|CLONE_SIGHAND|CLONE_THREAD|CLONE_SYSVSEM|CLONE_SETTLS|CLONE_PARENT_SETTID,
    parent_tidptr=0x7f5124a2b9d0, tls=0x7f5124a2b700, child_tidptr=0x7f5124a2b9d0) = 5106 <0.000042>
[pid 5106] 11:15:42.166585 set_robust_list(0x7f5124a2b9e0, 24 <unfinished ...>
[pid 5087] 11:15:42.166603 getpid( <unfinished ...>
[pid 5106] 11:15:42.166615 <... set_robust_list resumed> ) = 0 <0.000016>
[pid 5087] 11:15:42.166625 <... getpid resumed> ) = 5087 <0.000016>
[pid 5106] 11:15:42.166635 prctl(PR_SET_NAME, "ruby-timer-thr") = 0 <0.000009>
[pid 5106] 11:15:42.166669 poll([{fd=3, events=POLLIN}, {fd=5, events=POLLIN}], 2, -1 <unfinished ...>
[pid 5087] 11:15:42.167364 geteuid() = 1000 <0.000008>
[pid 5087] 11:15:42.167410 getegid() = 100 <0.000010>
[pid 5087] 11:15:42.167897 brk(0x2343000) = 0x2343000 <0.000014>
[pid 5087] 11:15:42.168282 getuid() = 1000 <0.000012>
[pid 5087] 11:15:42.168349 geteuid() = 1000 <0.000013>
[pid 5087] 11:15:42.168397 getgid() = 100 <0.000011>
[pid 5087] 11:15:42.168432 getegid() = 100 <0.000009>
[pid 5087] 11:15:42.168613 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000022>
[pid 5087] 11:15:42.168691
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/x86_64-linux/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000012>
[pid 5087] 11:15:42.168750 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000015>
[pid 5087] 11:15:42.168796
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000015>
[pid 5087] 11:15:42.168848
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/x86_64-linux/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000014>
```

# System Calls IX

```
[pid 5087] 11:15:42.168901 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000015>
[pid 5087] 11:15:42.168950 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000011>
[pid 5087] 11:15:42.168999
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/encdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = 8 <0.000016>
[pid 5087] 11:15:42.169049 fcntl(8, F_GETFD) = 0x1 (flags FD_CLOEXEC) <0.000008>
[pid 5087] 11:15:42.169081 fstat(8, {st_mode=S_IFREG|0755, st_size=90384, ...}) = 0 <0.000007>
[pid 5087] 11:15:42.169120 close(8) = 0 <0.000010>
[pid 5087] 11:15:42.169210 futex(0x7f5123e41048, FUTEX_WAKE_PRIVATE, 2147483647) = 0 <0.000008>
[pid 5087] 11:15:42.169247
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/encdb.so", O_RDONLY|O_CLOEXEC)
    = 8 <0.000014>
[pid 5087] 11:15:42.169283 read(8, "\177ELF\2\1\1\0\0\0\0\0\0\0\3\0\0\1\0\0\0\360\7\0\0\0\0\0"...,
    832) = 832 <0.000008>
[pid 5087] 11:15:42.169318 fstat(8, {st_mode=S_IFREG|0755, st_size=90384, ...}) = 0 <0.000007>
[pid 5087] 11:15:42.169353 mmap(NULL, 2109536, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 8, 0)
    = 0x7f5122f69000 <0.000015>
[pid 5087] 11:15:42.169392 mprotect(0x7f5122f6b000, 2097152, PROT_NONE) = 0 <0.000013>
[pid 5087] 11:15:42.169422 mmap(0x7f512316b000, 8192, PROT_READ|PROT_WRITE,
    MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 8, 0x2000) = 0x7f512316b000 <0.000018>
[pid 5087] 11:15:42.169480 close(8) = 0 <0.000008>
[pid 5087] 11:15:42.169547 mprotect(0x7f512316b000, 4096, PROT_READ) = 0 <0.000014>
[pid 5087] 11:15:42.170460
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000015>
[pid 5087] 11:15:42.170521
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/x86_64-linux/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000020>
```

# System Calls X

```
[pid 5087] 11:15:42.170586
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000012>
[pid 5087] 11:15:42.170636
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000015>
[pid 5087] 11:15:42.170688
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/x86_64-linux/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000012>
[pid 5087] 11:15:42.170740
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000012>
[pid 5087] 11:15:42.170787 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = -1 ENOENT (No such file or directory) <0.000012>
[pid 5087] 11:15:42.170833
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/trans/transdb.so",
    O_RDONLY|O_NONBLOCK|O_CLOEXEC) = 8 <0.000023>
[pid 5087] 11:15:42.170888 fstat(8, {st_mode=S_IFREG|0755, st_size=20200, ...}) = 0 <0.000010>
[pid 5087] 11:15:42.170930 close(8) = 0 <0.000009>
[pid 5087] 11:15:42.170992
    open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/trans/transdb.so",
    O_RDONLY|O_CLOEXEC) = 8 <0.000013>
[pid 5087] 11:15:42.171025 read(8, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\3\0\0\1\0\0\0\6\0\0\0\0\0...",
    832) = 832 <0.000009>
.....
.....
[pid 5087] 11:15:42.320297 write(1, "hello world", 11hello world) = 11 <0.000015>
[pid 5087] 11:15:42.320361 write(1, "\n", 1
) = 1 <0.000022>
.....
```

# System Calls XI

```
[pid 5106] 11:15:42.322634 exit(0)      = ?
[pid 5087] 11:15:42.322715 <... futex resumed> ) = 0 <0.000299>
[pid 5106] 11:15:42.322764 +++ exited with 0 +++
11:15:42.322787 munmap(0x7f51248ee000, 1052672) = 0 <0.000040>
11:15:42.322863 munmap(0x7f51249f7000, 200704) = 0 <0.000042>
11:15:42.322997 exit_group(0)        = ?
11:15:42.324387 +++ exited with 0 +++
```

#####

# Trace all 'open' syscalls

sudo strace -ttT -ff -e trace=open ruby -e 'puts "hello world"'

```
11:28:40.909716 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/x86_64/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.909820 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.909863 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/x86_64/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.909902 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = 3 <0.00001>
11:28:40.910057 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libpthread.so.0", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.910085 open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3 <0.000008>
11:28:40.910155 open("/usr/lib/libpthread.so.0", O_RDONLY|O_CLOEXEC) = 3 <0.000008>
11:28:40.910282 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libgmp.so.10", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.910307 open("/usr/lib/libgmp.so.10", O_RDONLY|O_CLOEXEC) = 3 <0.000007>
11:28:40.910418 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libdl.so.2", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.910445 open("/usr/lib/libdl.so.2", O_RDONLY|O_CLOEXEC) = 3 <0.000008>
11:28:40.910555 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libcrypt.so.1", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.910580 open("/usr/lib/libcrypt.so.1", O_RDONLY|O_CLOEXEC) = 3 <0.000008>
11:28:40.910702 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libm.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.910727 open("/usr/lib/libm.so.6", O_RDONLY|O_CLOEXEC) = 3 <0.000008>
11:28:40.910838 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
11:28:40.910863 open("/usr/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = 3 <0.000007>
11:28:40.911636 open("/usr/lib/locale/locale-archive", O_RDONLY|O_CLOEXEC) = 3 <0.000009>
```

# System Calls XII

```
11:28:40.911784 open("/proc/self/maps", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
```

```
strace: Process 3063 attached
```

```
[pid 3062] 11:28:40.921043 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921099 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/x86_64-linux/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921130 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921160 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921195 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/x86_64-linux/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921238 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921282 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921319 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.921471 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/encdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.922454 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/enc/trans/transdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.922493 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/x86_64-linux/enc/trans/transdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.922522 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/enc/trans/transdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.922551 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/enc/trans/transdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:40.922584 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/x86_64-linux/enc/trans/transdb.so", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
```

```
.....
.....
```

```
[pid 3062] 11:28:41.038858 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.039000 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.039980 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.040111 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.041014 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.041144 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.041765 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
[pid 3062] 11:28:41.041893 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1.0/lib/did_you_mean.rb", 0_RDONLY|0_CLOEXEC) = 3 <0.000017>
```

```
hello world
```

```
[pid 3063] 11:28:41.044118 +++ exited with 0 +++
```

```
11:28:41.045734 +++ exited with 0 +++
```

# System Calls XIII

#####

# Trace all file related syscalls

strace -ttT -ff -e trace=file ruby -e 'puts "hello world"'

```
11:31:25.733312 execve("/home/suresh/.rvm/rubies/ruby-2.4.1/bin/ruby", ["ruby", "-e", "puts \"hello world\""],
11:31:25.733617 access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory) <0.000011>
11:31:25.733668 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/x86_64/libruby.so.2.4", O_RDONLY|O_CLOEXEC) =
11:31:25.733712 stat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/x86_64", 0x7ffdabc866f0) = -1 ENOENT (No such
11:31:25.733758 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = -1 ENO
11:31:25.733795 stat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/tls", 0x7ffdabc866f0) = -1 ENOENT (No such file o
11:31:25.733831 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/x86_64/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = -1
11:31:25.733865 stat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/x86_64", 0x7ffdabc866f0) = -1 ENOENT (No such fil
11:31:25.733898 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libruby.so.2.4", O_RDONLY|O_CLOEXEC) = 3 <0.00001
11:31:25.734074 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libpthread.so.0", O_RDONLY|O_CLOEXEC) = -1 ENOENT
11:31:25.734109 open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3 <0.000010>
11:31:25.734198 open("/usr/lib/libpthread.so.0", O_RDONLY|O_CLOEXEC) = 3 <0.000010>
11:31:25.734343 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libgmp.so.10", O_RDONLY|O_CLOEXEC) = -1 ENOENT (N
11:31:25.734378 open("/usr/lib/libgmp.so.10", O_RDONLY|O_CLOEXEC) = 3 <0.000010>
11:31:25.734512 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libdl.so.2", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No
11:31:25.734550 open("/usr/lib/libdl.so.2", O_RDONLY|O_CLOEXEC) = 3 <0.000010>
11:31:25.734709 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libcrypt.so.1", O_RDONLY|O_CLOEXEC) = -1 ENOENT (
11:31:25.734746 open("/usr/lib/libcrypt.so.1", O_RDONLY|O_CLOEXEC) = 3 <0.000010>
11:31:25.734891 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libm.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No s
11:31:25.734926 open("/usr/lib/libm.so.6", O_RDONLY|O_CLOEXEC) = 3 <0.000009>
11:31:25.735063 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No s
11:31:25.735098 open("/usr/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = 3 <0.000009>
11:31:25.735888 open("/usr/lib/locale/locale-archive", O_RDONLY|O_CLOEXEC) = 3 <0.000011>
11:31:25.736051 open("/proc/self/maps", O_RDONLY|O_CLOEXEC) = 3 <0.000019>
strace: Process 9532 attached
[pid 9531] 11:31:25.745984 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site-ruby/2.4.0/enc/encdb.so", O
```

# System Calls XIV

```
[pid 9531] 11:31:25.746057 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/x86_64-linux/enc
[pid 9531] 11:31:25.746100 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/enc/encdb.so", 0_RDONLY
[pid 9531] 11:31:25.746142 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/enc/encdb.so",
[pid 9531] 11:31:25.746184 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/x86_64-linux/e
[pid 9531] 11:31:25.746225 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/enc/encdb.so", 0_RDONLY
[pid 9531] 11:31:25.746267 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/enc/encdb.so", 0_RDONLY|O_
[pid 9531] 11:31:25.746321 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/encdb.so"
[pid 9531] 11:31:25.746486 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/2.4.0/x86_64-linux/enc/encdb.so"
[pid 9531] 11:31:25.747511 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/enc/trans/transd
[pid 9531] 11:31:25.747556 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/2.4.0/x86_64-linux/enc
[pid 9531] 11:31:25.747596 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/site_ruby/enc/trans/transdb.so",
[pid 9531] 11:31:25.747641 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/enc/trans/trans
[pid 9531] 11:31:25.747687 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/vendor_ruby/2.4.0/x86_64-linux/e
.....
.....
[pid 9531] 11:31:25.888752 open("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.1
[pid 9531] 11:31:25.889119 lstat("/home", {st_mode=S_IFDIR|0755, st_size=4096, ...}) = 0 <0.000016>
[pid 9531] 11:31:25.889176 lstat("/home/suresh", {st_mode=S_IFDIR|0700, st_size=4096, ...}) = 0 <0.000011>
[pid 9531] 11:31:25.889213 lstat("/home/suresh/.rvm", {st_mode=S_IFDIR|0755, st_size=4096, ...}) = 0 <0.000011>
[pid 9531] 11:31:25.889252 lstat("/home/suresh/.rvm/rubies", {st_mode=S_IFDIR|0755, st_size=4096, ...}) = 0 <0
[pid 9531] 11:31:25.889293 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1", {st_mode=S_IFDIR|0755, st_size=4096,
[pid 9531] 11:31:25.889341 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib", {st_mode=S_IFDIR|0755, st_size=409
[pid 9531] 11:31:25.889389 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby", {st_mode=S_IFDIR|0755, st_siz
[pid 9531] 11:31:25.889436 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems", {st_mode=S_IFDIR|0755, s
[pid 9531] 11:31:25.889485 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0", {st_mode=S_IFDIR|0
[pid 9531] 11:31:25.889534 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems", {st_mode=S_IF
[pid 9531] 11:31:25.889589 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.
[pid 9531] 11:31:25.889637 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.
[pid 9531] 11:31:25.889690 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.
[pid 9531] 11:31:25.889739 lstat("/home/suresh/.rvm/rubies/ruby-2.4.1/lib/ruby/gems/2.4.0/gems/did_you_mean-1.
```



# System Calls XV

```
hello world
```

```
[pid 9532] 11:31:25.892159 +++ exited with 0 +++
```

```
11:31:25.893734 +++ exited with 0 +++
```

```
#####
```

```
# Or use 'perf trace'
```

```
perf trace ruby -e 'puts "hello world"'
```

```
0.024 ( 0.002 ms): ruby/23452 brk(                                ) = 0x1d430
0.042 ( 0.004 ms): ruby/23452 access(filename: 0xf05adc80, mode: R ) = -1 ENOE
0.050 ( 0.004 ms): ruby/23452 open(filename: 0x88f9ce10, flags: CLOEXEC ) = -1 ENOE
0.055 ( 0.002 ms): ruby/23452 stat(filename: 0x88f9ce10, statbuf: 0x7ffd88f9cee0 ) = -1 ENOE
0.058 ( 0.002 ms): ruby/23452 open(filename: 0x88f9ce10, flags: CLOEXEC ) = -1 ENOE
0.062 ( 0.002 ms): ruby/23452 stat(filename: 0x88f9ce10, statbuf: 0x7ffd88f9cee0 ) = -1 ENOE
0.065 ( 0.002 ms): ruby/23452 open(filename: 0x88f9ce10, flags: CLOEXEC ) = -1 ENOE
0.069 ( 0.002 ms): ruby/23452 stat(filename: 0x88f9ce10, statbuf: 0x7ffd88f9cee0 ) = -1 ENOE
0.072 ( 0.004 ms): ruby/23452 open(filename: 0x88f9ce10, flags: CLOEXEC ) = 3
0.077 ( 0.002 ms): ruby/23452 read(fd: 3, buf: 0x7ffd88f9d048, count: 832 ) = 832
0.081 ( 0.002 ms): ruby/23452 fstat(fd: 3</proc/23452/maps>, statbuf: 0x7ffd88f9cee0 ) = 0
.....
.....
114.789 ( 0.009 ms): ruby/23452 write(fd: 1</dev/pts/1>, buf: 0x1f38c20, count: 11 ) = 11
114.802 ( 0.003 ms): ruby/23452 write(fd: 1</dev/pts/1>, buf: 0x7f5cf02a53fc, count: 1 ) = 1
115.474 ( 0.041 ms): ruby/23452 futex(uaddr: 0x7f5cf07ae9d0, val: 23453, uaddr2: 0xca, val3: 140037148305872) =
    8.781 (106.695 ms): ruby-timer-thr/23453 ... [continued]: poll() = 2
115.519 (18446744073709.512 ms): ruby/23452 munmap(addr: 0x7f5cf0671000, len: 1052672
115.480 ( 0.002 ms): ruby-timer-thr/23453 read(fd: 3, buf: 0x7f5cf0587000, count: 1024
115.484 ( 0.001 ms): ruby-timer-thr/23453 read(fd: 5, buf: 0x7f5cf0587000, count: 1024
115.487 ( 0.002 ms): ruby-timer-thr/23453 close(fd: 3
115.495 ( 0.002 ms): ruby-timer-thr/23453 close(fd: 5
```

# System Calls XVI

```
115.502 ( 0.000 ms): ruby-timer-thr/23453 exit(  
115.519 ( 0.016 ms): ruby/23452 ... [continued]: munmap()) = 0  
115.539 ( 0.015 ms): ruby/23452 munmap(addr: 0x7f5cf077a000, len: 200704 ) = 0  
115.583 ( 0.000 ms): ruby/23452 exit_group( )
```

---

<sup>14</sup><http://duartes.org/gustavo/blog/post/system-calls/>

<sup>15</sup><http://www.brendangregg.com/blog/2014-05-11/strace-wow-much-syscall.html>

<sup>16</sup><http://man7.org/linux/man-pages/man2/ptrace.2.html>

# System Calls: Mode Switch vs Context Switch I

- Mode Switch: Switching from userspace to kernel (privilege 3 to privilege 0)
- Context Switch: Switching from process to another (== save all registers etc)
- Mode Switch != Context Switch
- A mode switch may / may not lead to a context switch, depending on the syscall/work involved
- Example: Making thousands of syscalls and not running into any context switches

# System Calls: Mode Switch vs Context Switch II

```
# Just make sure there are syscalls being made
strace -c -e getpid ruby -e 'for i in 0..10000; Process.pid; end'
% time      seconds  usecs/call   calls   errors syscall
-----
100.00      0.016122          2    10004         getpid
-----
100.00      0.016122          2    10004         total
```

```
# Trace the number of context switches
sudo perf stat -e context-switches ruby -e 'for i in 0..10000; Process.pid; end'
```

Performance counter stats for 'ruby -e for i in 0..10000; Process.pid; end':

```
                2          context-switches

0.097093505 seconds time elapsed
```

- Example: Make a http GET, observe how many context switches happen

# System Calls: Mode Switch vs Context Switch III

```
# How many context-switches we run into during a simple http GET call
sudo perf stat -e context-switches ruby -e "require 'net/http'; \
  Net::HTTP.get(URI.parse('https://www.google.co.in/'));"
```

```
Performance counter stats for 'ruby -e require 'net/http'; Net::HTTP.get(URI.parse('https://www.google.co.in/'))':

      17          context-switches

    0.336323461 seconds time elapsed
```

```
# Check where we are getting scheduled out
sudo perf record --call-graph dwarf -e context-switches \
  ruby -e "require 'net/http'; Net::HTTP.get(URI.parse('http://google.com/'));"
```

```
# Display the callstacks where it was scheduled out
sudo perf report --call-graph=flat,count --stdio --no-children
```

# Various process metrics

- CPU utilization
  - Remember: This also includes waiting to access memory
- Memory utilization
  - Heap
  - Stack
- Disk utilization

# Case study: Simple case - Curl I

```
$ strace -ff -ttT -e trace=network curl --silent -o /dev/null http://www.google.co.in
```

```
.....  
[pid 24286] 18:40:54.282834 socket(AF_INET, SOCK_DGRAM|SOCK_NONBLOCK, IPPROTO_IP) = 3 <0.000020>  
[pid 24286] 18:40:54.282889 connect(3, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("8.8.8.8")},  
.....  
[pid 24286] 18:40:54.322050 socket(AF_INET, SOCK_DGRAM, IPPROTO_IP) = 3 <0.000015>  
[pid 24286] 18:40:54.322088 connect(3, {sa_family=AF_INET, sin_port=htons(80), sin_addr=inet_addr("172.217.26.1  
[pid 24286] 18:40:54.322168 getsockname(3, {sa_family=AF_INET, sin_port=htons(38303), sin_addr=inet_addr("192.1  
[pid 24286] 18:40:54.322218 socket(AF_INET6, SOCK_DGRAM, IPPROTO_IP) = 3 <0.000014>  
[pid 24286] 18:40:54.322252 connect(3, {sa_family=AF_INET6, sin6_port=htons(80), inet_pton(AF_INET6, "2404:6800  
[pid 24286] 18:40:54.322451 +++ exited with 0 +++  
18:40:54.340534 socket(AF_INET, SOCK_STREAM, IPPROTO_TCP) = 3 <0.000052>  
.....  
18:40:54.340881 connect(3, {sa_family=AF_INET, sin_port=htons(80), sin_addr=inet_addr("172.217.26.163")}, 16) =  
.....  
18:40:54.344268 sendto(3, "GET / HTTP/1.1\r\nHost: www.google"... , 80, MSG_NOSIGNAL, NULL, 0) = 80 <0.000022>  
18:40:54.474244 recvfrom(3, "HTTP/1.1 200 OK\r\nDate: Tue, 08 A"... , 102400, 0, NULL, NULL) = 13880 <0.000012>  
18:40:54.477151 recvfrom(3, "umpException=window._.DumpExcep"... , 102400, 0, NULL, NULL) = 1275 <0.000008>  
18:40:54.478162 +++ exited with 0 +++
```

# Case study: How Ruby threading works I

- MRI Ruby is single threaded
- Even though it supports multi-threading, only one thread can execute ruby code at a given time.
- Why: MRI Ruby VM is not multi-thread safe: A global lock must be acquired by a thread before it can execute ruby code. It is called GVL (Global VM Lock) in Ruby (like GIL in python).
- So in a Ruby process, at any given time, only one thread can execute for all practical purposes, everybody else must wait for the currently running thread to release the GVL lock.
- Now there are two scenarios w.r.t multi-threading
  - 1 Current thread is voluntarily giving up the lock (by calling blocking functions, like sleep, or waiting on network call etc).
  - 2 Current thread doesn't call any blocking function, just busily does some work (like big regex match, or "loop do end" etc).



# Case study: How Ruby threading works II

For case 1, this is the most straightforward, RubyVM doesn't have to do anything. Here is an example, sleep function <sup>17</sup>, notice that:

- It releases the GVL lock
  - Goes to sleep
  - As soon as it wakes up, it will try to acquire the GVL lock, only then it can proceed further. If it can't acquire GVL lock immediately, it will have to wait. For case 2, we will see below how it works.
- How Ruby scheduler works:
    - If a thread acquires GVL, if it continues to run without calling any blocking function, it can unfairly starve all other threads (all those threads will be waiting to acquire GVL, which this thread doesn't release).
    - So Ruby implements a sort of co-operative scheduler: a pseudo co-operative interrupt mechanism



# Case study: How Ruby threading works III

- A timer thread is started, it keeps running forever. Periodically, it will set a "flag" in the current running thread to indicate it should stop running and give up the GVL. The function that actually does it:  
[https://github.com/ruby/ruby/blob/ruby\\_2\\_2/thread.c#L3817-L3834](https://github.com/ruby/ruby/blob/ruby_2_2/thread.c#L3817-L3834)
- Now the current thread that is running will check this flag periodically to see if it needs to stop running. If so, it will give up the lock and go back to wait list.

---

<sup>17</sup>[https://github.com/ruby/ruby/blob/ruby\\_2\\_2/thread.c#L3817-L3834](https://github.com/ruby/ruby/blob/ruby_2_2/thread.c#L3817-L3834)

# Case study: How NodeJS eventing/threading works I



- Node.js has single thread + event loop architecture
- It uses libuv for "handling" events
- How does it handle blocking syscalls?
  - It uses non-blocking IO wherever possible (if OS supports) and fallbacks to threadpool for blocking IO
    - Non-blocking: TCP, UDP, Pipes etc
    - Blocking: DNS lookups, Disk read/writes etc
- Let's analyze it from the PoV of OS to validate this

# Case study: How NodeJS eventing/threading works II

```
# strace -ff -ttT -e trace=network,write,read node -p 'console.log('##### Process ID #####: ${process.pid}'); \
req=http.get("http://www.google.co.in", function(res) { var body=""; \
res.on("data",function(data){body+=data;}); \
res.on("end", function() {console.log("=====")} )});console.log("Done")'
.....
[pid 12551] 15:42:26.773798 write(9, "##### Process ID #####...", 41##### Process ID #####
.....
[pid 12551] 15:42:26.799280 write(9, "Done\n", 5) = 5 <0.000018>
.....
[pid 12557] 15:42:26.799574 socket(AF_INET, SOCK_DGRAM|SOCK_NONBLOCK, IPPROTO_IP) = 12 <0.000014>
[pid 12557] 15:42:26.799615 connect(12, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("8.8.8.8")},
[pid 12557] 15:42:26.799675 sendmmsg(12, [{msg_hdr={msg_name=NULL, msg_namelen=0, msg_iov=[{iov_base="\377F\10
[pid 12557] 15:42:26.838085 recvfrom(12, "\377F\201\200\0\1\0\1\0\0\0\3www\6google\2co\2in\0\0\1"... , 2048, 0
.....
[pid 12557] 15:42:26.838542 socket(AF_INET, SOCK_DGRAM, IPPROTO_IP) = 12 <0.000014>
[pid 12557] 15:42:26.838577 connect(12, {sa_family=AF_INET, sin_port=htons(0), sin_addr=inet_addr("216.58.196.9
[pid 12557] 15:42:26.838618 getsockname(12, {sa_family=AF_INET, sin_port=htons(56665), sin_addr=inet_addr("192.
[pid 12557] 15:42:26.838666 socket(AF_INET6, SOCK_DGRAM, IPPROTO_IP) = 12 <0.000013>
[pid 12557] 15:42:26.838713 connect(12, {sa_family=AF_INET6, sin6_port=htons(0), inet_pton(AF_INET6, "2404:6800
.....
[pid 12551] 15:42:26.839723 socket(AF_INET, SOCK_STREAM|SOCK_CLOEXEC|SOCK_NONBLOCK, IPPROTO_IP) = 12 <0.000052>
[pid 12551] 15:42:26.839811 connect(12, {sa_family=AF_INET, sin_port=htons(80), sin_addr=inet_addr("216.58.196.
[pid 12551] 15:42:26.842964 getsockopt(12, SOL_SOCKET, SO_ERROR, [0], [4]) = 0 <0.000010>
[pid 12551] 15:42:26.843780 write(12, "GET / HTTP/1.1\r\nHost: www.google"... , 61) = 61 <0.000038>
[pid 12551] 15:42:26.968154 read(12, "HTTP/1.1 200 OK\r\nDate: Mon, 07 A"... , 65536) = 13880 <0.000024>
[pid 12551] 15:42:26.972625 read(12, "f1f1;background-image:-ms-linea"... , 65536) = 27760 <0.000025>
[pid 12551] 15:42:26.974256 read(12, "<a class=gbmt href=\"http://www.g"... , 65536) = 7844 <0.000016>
[pid 12551] 15:42:26.974416 read(12, "", 65536) = 0 <0.000011>
[pid 12551] 15:42:26.976656 write(9, "=====\n", 23) = 23 <0.000027>
.....
```

# Case study: How NodeJS eventing/threading works III

## ● Observations

- We do see the DNS look up being made from a separate thread
- Making http request(= TCP), writing to console(= TTY) etc are being made from main thread
- But we also see multiple 'connect' calls to the server, why and where do they originate from?

```
sudo perf trace -T -e connect --call-graph=dwarf node -p 'console.log('##### Process ID ####:${process.pid}');\n  req=http.get("http://www.google.co.in", function(res) { var body=""; \n  res.on("data",function(data){body+=data;}); \n  res.on("end", function() {console.log("=====")} )});console.log("Done")'\n##### Process ID #####: 10113\nDone\n160498846.387 ( 0.026 ms): :10119/10119 connect(fd: 13, uservaddr: 0x7fbce3e345a0, addrln: 110\n                                [0] ([unknown])\n160498846.432 ( 0.004 ms): :10119/10119 connect(fd: 12<socket:[32719970]>, uservaddr: 0x7fbce3e35660, addrln:\n                                [0] ([unknown])\n160498847.194 ( 0.016 ms): :10119/10119 connect(fd: 12<socket:[32719970]>, uservaddr: 0x7fbcdc002e08, addrln:\n                                [0] ([unknown])\n160498848.516 ( 0.014 ms): node/10119 connect(fd: 12<socket:[32719970]>, uservaddr: 0x7fbce3e36dcc, addrln: 16\n                                __GI___libc_connect (/usr/lib/libpthread-2.25.so)\n                                [0xffff80431ec83d32] (/usr/lib/libresolv-2.25.so)\n                                [0xffff80431ec84c1e] (/usr/lib/libresolv-2.25.so)\n                                __libc_res_nquery (/usr/lib/libresolv-2.25.so)\n                                [0xffff80431ec833af] (/usr/lib/libresolv-2.25.so)\n                                __libc_res_nsearch (/usr/lib/libresolv-2.25.so)
```

# Case study: How NodeJS eventing/threading works IV

```
_nss_dns_gethostbyname4_r (/usr/lib/libnss_dns-2.25.so)
gai_inet.constprop.5 (/usr/lib/libc-2.25.so)
getaddrinfo (/usr/lib/libc-2.25.so)
[0xffff8043161249e9] (/usr/lib/libuv.so.1.0.0)
[0xffff80431611ca54] (/usr/lib/libuv.so.1.0.0)
start_thread (/usr/lib/libpthread-2.25.so)
__clone (/usr/lib/libc-2.25.so)
160498887.935 ( 0.014 ms): node/10119 connect(fd: 12<socket:[32719970]>, servaddr: 0x7fbcdc001fb0, addrlen: 16
__GI___libc_connect (/usr/lib/libc-2.25.so)
getaddrinfo (/usr/lib/libc-2.25.so)
160498887.964 ( 0.018 ms): node/10119 connect(fd: 12<socket:[32719970]>, servaddr: 0x7fbcdc002000, addrlen: 28
__GI___libc_connect (/usr/lib/libc-2.25.so)
getaddrinfo (/usr/lib/libc-2.25.so)
Failed to open /tmp/perf-10113.map, continuing without symbols
160498889.020 ( 0.058 ms): node/10113 connect(fd: 12<socket:[32718959]>, servaddr: 0x7ffebcc03020, addrlen: 16
__GI___libc_connect (/usr/lib/libpthread-2.25.so)
uv_tcp_connect (/usr/lib/libuv.so.1.0.0)
node::TCPWrap::Connect (/usr/bin/node)
v8::internal::FunctionCallArguments::Call (/usr/bin/node)
[0xf0b8] (/usr/bin/node)
[0xf450] (/usr/bin/node)
v8::internal::Builtin_HandleApiCall (/usr/bin/node)
[0x3dbac15043a7] (/tmp/perf-10113.map)
```

# Case study: How Golang threading works I

TODO

- 1 System
- 2 Application
- 3 Profilers**
- 4 Debuggers
- 5 Books



- Sampling vs Precise
  - Tail latency
- Use cases
  - CPU intensive
    - CPU cache usage and Cache contention
  - Blocking (off-cpu)
  - Language specific
- perf
- stackprof

- Core vs HyperThreading

<https://github.com/andikleen/pmu-tools> <https://github.com/andikleen/pmu-tools/wiki/toplev-manual>

```
# Run it on the same core (HyperThreading)
echo -n "2 6" | xargs -I'{}' --delimit ' ' --max-args=1 --max-procs=$(nproc) \
  sudo sh -c 'chrt -f 99 dd if=/dev/zero bs=8M count=128 | \
  toplev.py --quiet --single-thread -l3 taskset -c {} chrt -f 99 gzip > /dev/null'
128+0 records in
128+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 14.2675 s, 75.3 MB/s
BE          Backend_Bound:          60.77 % Slots [ 6.25%]
BE/Mem      Backend_Bound.Memory_Bound: 44.44 % Slots [ 6.25%]
BE/Mem      Backend_Bound.Memory_Bound.L1_Bound: 25.39 % Clocks [ 6.26%] BN

128+0 records in
128+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 14.3168 s, 75.0 MB/s
BE          Backend_Bound:          51.97 % Slots [ 6.25%]
BE/Mem      Backend_Bound.Memory_Bound: 43.18 % Slots [ 6.25%]
BE/Mem      Backend_Bound.Memory_Bound.L1_Bound: 46.23 % Clocks [ 6.26%] BN

# Run it on different cores
echo -n "2 3" | xargs -I'{}' --delimit ' ' --max-args=1 --max-procs=$(nproc) \
  sudo sh -c 'chrt -f 99 dd if=/dev/zero bs=8M count=128 | \
```

```
toplev.py --quiet --single-thread -l3 taskset -c {} chrt -f 99 gzip > /dev/null'
```

```
128+0 records in
128+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 9.90494 s, 108 MB/s
128+0 records in
128+0 records out
1073741824 bytes (1.1 GB, 1.0 GiB) copied, 9.90628 s, 108 MB/s
BE          Backend_Bound: 42.88 % Slots [ 6.29%] BN
BE/Mem      Backend_Bound.Memory_Bound: 21.33 % Slots [ 6.22%]
BE/Core     Backend_Bound.Core_Bound: 21.55 % Slots [ 6.22%]
BE/Mem      Backend_Bound.Memory_Bound.L1_Bound: 32.82 % Clocks [ 6.20%]
BE/Core     Backend_Bound.Core_Bound.Ports_Utilization: 36.13 % Clocks [ 6.26%]
BE          Backend_Bound: 42.66 % Slots [ 6.25%] BN
BE/Mem      Backend_Bound.Memory_Bound: 20.95 % Slots [ 6.25%]
BE/Core     Backend_Bound.Core_Bound: 21.71 % Slots [ 6.25%]
BE/Mem      Backend_Bound.Memory_Bound.L1_Bound: 32.87 % Clocks [ 6.22%]
BE/Core     Backend_Bound.Core_Bound.Ports_Utilization: 36.66 % Clocks [ 6.18%]
```

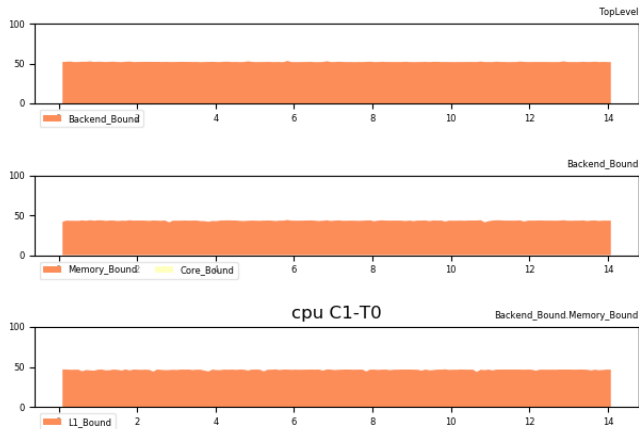


Figure: Performance on shared core

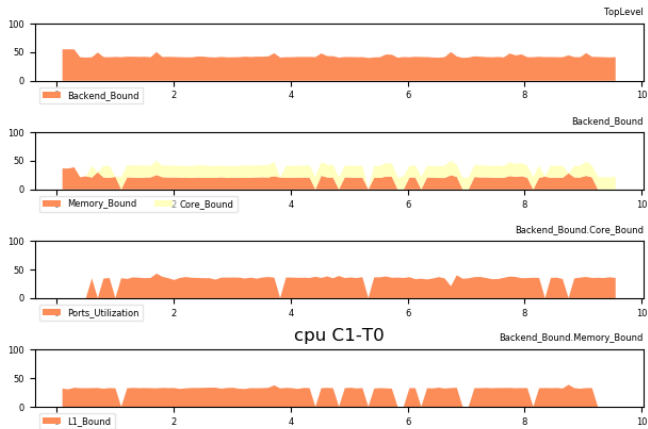


Figure: Performance on dedicated core

- 1 System
- 2 Application
- 3 Profilers
- 4 Debuggers**
- 5 Books

# Debuggers

- gdb
- rbtrace

- 1 System
- 2 Application
- 3 Profilers
- 4 Debuggers
- 5 Books**



# Books for reference I

- How computers work

- Code: The Hidden Language of Computer Hardware and Software
- Hardware
  - The Indispensable PC Hardware Book Hardcover by Mr Hans-Peter Messmer
  - Modern Processor Design: Fundamentals of Superscalar Processors by John Paul Shen
  - Pentium Pro and Pentium II System Architecture (2nd Edition) by Tom Shanley

- OS

- Operating Systems: Three Easy Pieces -  
<http://pages.cs.wisc.edu/~remzi/OSTEP/>
- xv6 -  
<https://pdos.csail.mit.edu/6.828/2014/xv6/book-rev8.pdf>

- Linux

- Linux Kernel Development by Robert Love

- The Linux Programming Interface – A Linux and UNIX System Programming Handbook by Michael Kerrisk
- Linux System Programming (2 edition) by Robert Love
- Understanding The Linux Network Internals by Benvenuti
- Performance related
  - Site Reliability Engineering: How Google Runs Production Systems by Niall Murphy, Jennifer Petoff, Chris Jones
  - The Art of Computer Systems Performance Analysis: Techniques for Experimental Design, Measurement, Simulation and Modeling by Raj Jain
  - Systems Performance: Enterprise and the Cloud by Brendan Gregg